

MEMCALIB: Benchmarking and Optimizing Memory Use in LLM Agents

Ruike Cao^{1,2,*}, Fanyu Zhao^{2,3,*}, Fugen Yao^{2,†}, Liang Dong², Jian Xu², Guanjin Jiang², Yifei Zhao³, Han Zhang³ and Li Xiao^{1,†}

¹University of Science and Technology of China, ²Qwen Business Unit of Alibaba, ³Fudan University

*Equal contribution. †Corresponding authors.

The effectiveness of agent memory ultimately depends on whether the underlying LLM gives each memory in context an appropriate degree of influence over its response. Yet this capability has remained largely overlooked. To assess this capability, we introduce MEMCALIB, a benchmark grounded in realistic memory-system scenarios for evaluating memory use and advancing optimization algorithms. Results on the MEMCALIB test set reveal that frontier open- and closed-source models struggle to use memory appropriately. They frequently over-use or under-use memory rather than matching each proposition’s actual use to its target level, leading to biased, low-quality responses. Experiments with common post-training algorithms, including group relative policy optimization and on-policy self-distillation, further reveal a clear directional skew: trained models improve in one direction while deteriorating in the other. We therefore propose MEMCALIB-RL, an ordered bidirectional counterfactual credit-assignment algorithm that separates over- and under-use signals and localizes their credit to response tokens through exact atom ablation. Results across model families and scales (Qwen3-8B, Ministral-3-8B-Instruct, and Qwen3.5-35B-A3B) show that MEMCALIB-RL achieves the best overall performance while better balancing over-use and under-use, with gains generalizing beyond MEMCALIB in external benchmark evaluation. Further experiments support its design choices and robustness and provide insight into its training dynamics.¹

1. Introduction

Recent advances in reasoning, planning, and coding capabilities, together with the integration of memory and tools, have transformed LLMs from static conditional generators into adaptive policies that interact with external environments [37]. This transformation has given rise to LLM-based agents, in which memory serves as a core capability supporting long-horizon reasoning, continual adaptation, and coherent interaction with complex environments [8, 19]. Although memory systems may differ in how they form, update, retrieve, and filter memories, their effectiveness ultimately depends on whether the LLM gives each memory in context the appropriate level of influence. Post-retrieval filtering cannot guarantee a noise-free context: some supplied memories may remain irrelevant or outdated, while even useful memories differ in the scope of influence they should have. Moreover, real memory systems often return summaries, profiles, or trajectories in which a single memory block contains multiple atomic propositions [41, 22, 15]. Within the same block, some atoms may be noise, some should provide only local support, and others should control a core conclusion. The model must therefore make a finer decision than use or ignore: it must determine, at the atomic level, how each proposition should shape its response. For each atomic proposition in the supplied memory context, its appropriate level of influence on the response is determined based on the current task: the proposition should leave no answer-specific footprint (IGNORE), provide bounded local support (BOUND), or control a material conclusion (CONTROL). Influence that exceeds or falls below the target level is harmful: over-use can let irrelevant or outdated memories override current evidence, distort core conclusions, or cause over-personalization [9, 34]; under-use can discard preferences, constraints, and experience that should govern the answer [4, 38, 35]. Given that the effectiveness of agent memory systems ultimately depends on the LLM’s ability to use memory appropriately, a fundamental question remains largely overlooked:

¹Code and data: <https://github.com/Quark-Medical/memcalib>

[†]Corresponding authors: Fugen Yao (fugen.yfg@alibaba-inc.com) and Li Xiao (xiaoli11@ustc.edu.cn).




Figure 1 | Memory use and optimization. (a) An example of LLM agents over-using irrelevant memory and under-using relevant constraints. (b) A schematic comparison of optimization behavior: GRPO and OPSD can reduce both errors in principle but tend to prioritize one objective in practice due to coarse credit assignment, while MEMCALIB-RL better balances the two for better overall performance.

Are LLMs really capable of using memory appropriately to shape their responses?

Recent personalization benchmarks have already revealed limitations in LLM-based personalized assistants: they may over- or under-use personalized preference memories, causing them to over-personalize, become sycophantic, or fail to apply relevant preferences [9, 34, 4, 38] (see Appendix A for further discussion). Taken together, these limitations motivate us to re-examine LLMs’ ability to use memory appropriately. We therefore introduce MEMCALIB, a benchmark spanning health, general assistance, and coding, comprising 15,000 examples: a 13,500-example training set to facilitate innovation in optimization algorithms and a disjoint 1,500-example test set for evaluation. MEMCALIB captures a realistic challenge faced by LLMs in agent systems: given a query and a variable number of natural composite memory blocks, each of which contains a variable number of atomic propositions, a model must calibrate each proposition’s influence on its response. For evaluation, an LLM-based judge [40] applies atom-specific rubrics to classify the model’s actual use of each atom in its response as IGNORE, BOUND, or CONTROL. Comparing the target and actual use levels then quantifies over-use, under-use, and overall memory-use performance.

Results on the test set reveal widespread mismatches between atoms’ actual and target use levels across frontier open- and closed-source models. Most models also exhibit clear *directional skew*: they perform relatively well with respect to one error type but worse on the other, indicating a globally aggressive or conservative memory-use policy rather than calibrating each proposition to the current query. Moreover, experiments show that widely used post-training algorithms, including group relative policy optimization (GRPO) [29] and on-policy self-distillation (OPSD) [39], also exhibit clear directional skew: the trained models improve in one direction while deteriorating in the other, a phenomenon which we call the *calibration seesaw*. We therefore propose MEMCALIB-RL, an ordered bidirectional counterfactual credit-assignment method that decomposes the matrix of target and actual use levels into nine reward channels and then uses exact atom ablation and target-aware counterfactual likelihood differences to localize each channel’s credit to response tokens. Across model families and scales (Qwen3-8B, Ministral-3-8B-Instruct, and Qwen3.5-35B-A3B), MEMCALIB-RL delivers the best overall performance while better balancing over-use and under-use. Figure 1 summarizes the memory-use challenge and the contrasting optimization directions. Furthermore, these gains extend beyond MEMCALIB to the external RPEval benchmark [4]. Ablation and sensitivity analyses support our localization design. Mechanism analyses reveal how credit allocation evolves during training, while evaluation with an alternative Judge supports the robustness of the overall method comparison. Our contributions are:

- We identify a long-overlooked problem: LLMs often fail to use memory appropriately to shape their responses and exhibit clear directional skew between over-use and under-use.
- We introduce MEMCALIB, a benchmark that evaluates whether models match each proposition’s influence to its target level within natural composite memory blocks across health, general assistance, and coding.
- We propose MEMCALIB-RL, an ordered bidirectional counterfactual credit-assignment algorithm. Ex-

tensive experiments across model families, scales, and benchmarks demonstrate consistent gains and a better balance between over-use and under-use.

2. Benchmarking Memory Use in LLMs

2.1. Problem formulation

Consider an LLM-based agent that receives a user query together with memory blocks retrieved from its memory system. The LLM must determine how each block’s constituent propositions should shape its response. Let x denote the query, $M = \{b_n\}_{n=1}^N$ the supplied memory blocks, and π_θ the LLM policy with parameters θ . The model generates $y \sim \pi_\theta(\cdot | x, M)$. Each block b_n comprises K_n atomic propositions, $b_n = \{m_{n,k}\}_{k=1}^{K_n}$. For the current query x , each atomic proposition $m_{n,k}$ has an ideal use level $\ell_{n,k} \in \{\text{IGNORE}, \text{BOUND}, \text{CONTROL}\}$, where **IGNORE** leaves no answer-specific footprint, **BOUND** provides bounded local support, and **CONTROL** determines a material conclusion, constraint, or recommendation. An LLM that uses memory appropriately should let each atom influence its response at that atom’s ideal level.

2.2. Benchmark Construction

We therefore construct MEMCALIB from eight public question-answering datasets spanning health, general assistance, and coding to evaluate whether LLMs can use memory appropriately across a broad range of scenarios. Our core construction strategy is to separate each source sample into a current query and memory atoms containing information not included in the query, and annotate the extracted atoms with ideal use levels. The source context generally supports the current request, so most extracted atoms are labeled **BOUND** or **CONTROL**; we therefore add controlled **IGNORE** distractor atoms. We then generate atom-specific rubrics for all atoms, specifying response-text criteria for assessing their actual use levels. The atoms are then assembled into composite blocks and rewritten as fluent memory text while preserving every proposition. To ensure data quality, we use a six-stage LLM-based construction pipeline with deterministic checks and two independent LLM-based semantic reviews, which is iteratively refined based on human review of freshly constructed pilot batches and then frozen for full-scale construction (see Appendix B.1 for details). Using this pipeline, we construct 15,000 MEMCALIB examples spanning health, general assistance, and coding, with a stratified partition yielding 13,500 training and 1,500 disjoint test examples. Detailed benchmark statistics are provided in Appendix B.2.

2.3. Evaluation protocol and metrics

We adopt rubric-guided LLM-as-a-Judge evaluation [40, 11, 14] to assess the actual use level $\hat{\ell}_{n,k}$ of each atom $m_{n,k}$ in response y as **IGNORE**, **BOUND**, or **CONTROL**. The judge prompt is provided in Appendix E.2. To quantify deviations from the ideal levels, we define rank to map **IGNORE**, **BOUND**, and **CONTROL** to 0, 1, 2, respectively. We then define the response-level over-use and under-use totals by summing how far each atom’s actual use level exceeds or falls below its ideal level, respectively:

$$O = \sum_{n,k} \max(\text{rank}(\hat{\ell}_{n,k}) - \text{rank}(\ell_{n,k}), 0), \quad U = \sum_{n,k} \max(\text{rank}(\ell_{n,k}) - \text{rank}(\hat{\ell}_{n,k}), 0).$$

Let O_j and U_j denote the over-use and under-use totals for the j -th response in evaluation set $\mathcal{E}$, respectively. With a geometric decay factor $\rho = 0.5$ (each exponential term halves per unit increase in its exponent), we define sample-level Memory Overuse Severity (sMOS) and sample-level Memory Underuse Severity (sMUS) as:

$$\text{sMOS} = \frac{1}{|\mathcal{E}|} \sum_{j=1}^{|\mathcal{E}|} (1 - \rho^{O_j}), \quad \text{sMUS} = \frac{1}{|\mathcal{E}|} \sum_{j=1}^{|\mathcal{E}|} (1 - \rho^{U_j}).$$

We also report the Any-Overuse Rate (AOR) and Any-Underuse Rate (AUR) as the proportions of responses with at least one over-use or under-use error, respectively. To assess overall memory-use performance, we define

Table 1 | Evaluation results of representative models on the MEMCALIB test set (mean $\pm$ standard deviation over three seeds). Red and blue denote the best and second-best result in each column.

Model	SCS $\uparrow$	Exact $\uparrow$	sMOS $\downarrow$	sMUS $\downarrow$	AOR $\downarrow$	AUR $\downarrow$
GPT-5.6-SOL	46.25$\pm$0.60	28.40$\pm$0.85	38.45$\pm$0.15	22.92 $\pm$ 0.40	53.33$\pm$0.45	35.82 $\pm$ 0.51
Claude Sonnet 4.6	36.44$\pm$0.34	19.09 $\pm$ 0.33	48.89 $\pm$ 0.62	24.77 $\pm$ 0.37	65.40 $\pm$ 0.77	37.93 $\pm$ 0.38
Kimi-K2.6	35.54 $\pm$ 0.18	20.00$\pm$0.28	53.83 $\pm$ 0.19	19.94$\pm$0.60	68.67 $\pm$ 0.63	31.56$\pm$0.98
Gemini 3.5 Flash	34.96 $\pm$ 0.53	17.96 $\pm$ 0.58	51.65 $\pm$ 0.45	24.73 $\pm$ 0.14	68.27 $\pm$ 0.38	37.69 $\pm$ 0.25
GLM-5.2	34.40 $\pm$ 0.45	18.20 $\pm$ 0.61	52.13 $\pm$ 0.42	23.42 $\pm$ 0.31	67.53 $\pm$ 0.88	35.73 $\pm$ 0.71
Qwen3.8-Max	34.22 $\pm$ 0.52	17.80 $\pm$ 0.48	50.11 $\pm$ 0.33	27.46 $\pm$ 0.32	66.09 $\pm$ 0.51	41.16 $\pm$ 0.21
DeepSeek-V4-Flash	33.72 $\pm$ 0.76	16.96 $\pm$ 0.93	55.64 $\pm$ 1.08	20.86 $\pm$ 0.59	73.31 $\pm$ 1.01	32.64 $\pm$ 0.76
Qwen3-8B	31.17 $\pm$ 0.21	15.29 $\pm$ 0.23	37.19$\pm$0.62	45.66 $\pm$ 0.36	50.24$\pm$0.87	61.60 $\pm$ 0.24
Qwen3.5-35B-A3B	26.54 $\pm$ 0.28	12.24 $\pm$ 0.25	64.76 $\pm$ 0.48	19.93$\pm$0.32	80.16 $\pm$ 0.83	31.11$\pm$0.60

the Sample Calibration Score (SCS) and Exact Calibration (Exact) accounting for both over-use and under-use:

$$\text{SCS} = \frac{1}{|\mathcal{E}|} \sum_{j=1}^{|\mathcal{E}|} \rho^{O_j+U_j}, \quad \text{Exact} = \frac{1}{|\mathcal{E}|} \sum_{j=1}^{|\mathcal{E}|} \mathbf{1}[O_j + U_j = 0].$$

Higher SCS and Exact scores, together with lower error severity and rates, indicate more appropriate memory use. All six metrics are defined on $[0, 1]$ and rescaled to $[0, 100]$ for reporting.

2.4. Memory-Use Performance Across Models

Having established atom-specific rubrics and comprehensive evaluation metrics, we evaluate representative open- and closed-source models on the 1,500-example MEMCALIB test set, using DeepSeek-V4-Pro [2] to assess each atom’s actual use level. We run each model in non-thinking mode using three independent seeds, with both temperature and top- p set to 1.

The evaluation results in Table 1 show that all evaluated models struggle to use memory appropriately, with low SCS and Exact scores even among frontier closed-source models. Moreover, larger models do not necessarily use memory more appropriately: Qwen3-8B outperforms Qwen3.5-35B-A3B in both SCS and Exact. Beyond SCS and Exact, the directional metrics reveal a clear skew among the evaluated LLMs: all except Qwen3-8B over-use memory more severely and more frequently than they under-use it, whereas Qwen3-8B shows the reverse. GPT-5.6-SOL achieves the strongest overall performance because its errors are comparatively low in both directions. This outcome also illustrates that SCS and Exact are comprehensive metrics that take both over-use and under-use into account: strength in one direction cannot compensate for substantial errors in the other. This directional skew further suggests that models rely on a broad prior over whether memory should be trusted, leaving individual propositions poorly calibrated to the current query. To verify the reliability of the use-level judgments underlying these results, we have a human annotator label the actual use levels of sampled atoms. The results show 96.7% agreement between the Judge and human annotations on 150 naturally sampled judgments (95% Wilson CI: 92.4%–98.6%; Cohen’s $\kappa = 0.872$), with disagreements concentrated at the BOUND/CONTROL boundary (Appendix C).

3. MEMCALIB-RL: Ordered Bidirectional Credit Assignment

A single response can use some memory atoms appropriately while over- or under-using others, yet standard GRPO collapses these outcomes into a single advantage shared by every token [29]. In our experiments (Section 4), this coarse supervision yields a calibration seesaw: reducing one error type increases the other. We therefore propose MEMCALIB-RL (Figure 2), which separates the nine ideal–actual use transitions into distinct reward channels and uses bidirectional counterfactual evidence to redistribute each channel’s advantage across response tokens while preserving its response-level mean.




Figure 2 | Overview of MEMCALIB-RL. (a) Each atom’s ideal and actual use levels determine its assignment to one of nine reward channels. (b) Exact atom ablation yields token-level log-likelihood differences for the same response under full and ablated memories. (c) These signals guide channel-specific advantage redistribution toward supported or suppressed tokens while preserving each channel’s response-level mean.

3.1. Ordered reward channels

For each query–memory pair (x, M) , we sample a rollout group $\{y_i\}_{i=1}^G$, with $y_i \sim \pi_\theta(\cdot | x, M)$. Each atom’s ideal use level $\ell_{n,k}$ is fixed across the group, whereas its actual use level $\hat{\ell}_{n,k}(y_i)$ is assessed separately for each response y_i . For each y_i , we categorize the atoms according to $(\ell_{n,k}, \hat{\ell}_{n,k}(y_i))$ into nine ideal–actual reward channels indexed by h . To simplify notation, we denote IGNORE, BOUND, and CONTROL by A , B , and C , respectively, and accordingly define a symbolic label for each of the nine channels, as summarized in Table 8. Let $\ell^{(h)}$ and $\hat{\ell}^{(h)}$ denote the fixed ideal and actual levels defining channel h , and let $\mathcal{I}_i^{(h)}$ denote its atom-index set for response y_i :

$$\mathcal{I}_i^{(h)} = \{(n, k) : \ell_{n,k} = \ell^{(h)}, \hat{\ell}_{n,k}(y_i) = \hat{\ell}^{(h)}\}.$$

To reward correct atom-level use and penalize incorrect use, we define each channel reward as

$$R_i^{(h)} = \frac{|\mathcal{I}_i^{(h)}|}{Z^{(h)}} \begin{cases} +1, & \ell^{(h)} = \hat{\ell}^{(h)}, \\ -|\text{rank}(\hat{\ell}^{(h)}) - \text{rank}(\ell^{(h)})|, & \ell^{(h)} \neq \hat{\ell}^{(h)}, \end{cases}$$

$$Z^{(h)} = \max\left(1, \sum_{n=1}^N \sum_{k=1}^{K_n} \mathbf{1}[\ell_{n,k} = \ell^{(h)}]\right).$$

3.2. Bidirectional counterfactual localization

Given the channel rewards, GRPO aggregates them before group normalization, whereas GDPO preserves channel-specific signals by normalizing each channel separately before aggregation [18]. Both still assign the resulting response-level advantage uniformly to all tokens and therefore cannot localize each channel’s credit to the relevant tokens (see Appendix D.2 for details on both methods). In MEMCALIB-RL, we use fixed-response counterfactual log-likelihood differences to localize the influence of each channel’s memory atoms and guide advantage redistribution.

Specifically, for response y_i , we separately localize each channel h whose atom set satisfies $\mathcal{I}_i^{(h)} \neq \emptyset$, except (IGNORE, IGNORE), which is excluded because correctly ignored atoms leave no atom-specific response footprint. To localize channel h , we remove the atoms indexed by $\mathcal{I}_i^{(h)}$, leaving all other atoms and their order unchanged to obtain $\tilde{M}_i^{(h)}$. We then perform teacher-forcing inference on the same generated response under $\tilde{M}_i^{(h)}$ using the same rollout policy and define the log-likelihood difference for its t -th token $y_{i,t}$, conditioned on the preceding tokens $y_{i,<t}$, as

$$d_{i,t}^{(h)} = \log p_\theta(y_{i,t} | x, M, y_{i,<t}) - \log p_\theta(y_{i,t} | x, \tilde{M}_i^{(h)}, y_{i,<t}).$$

Since the response is fixed, $d_{i,t}^{(h)} > 0$ identifies tokens supported by the removed atoms, whereas $d_{i,t}^{(h)} < 0$ identifies tokens suppressed by them. We construct each channel’s localization signal according to the following

rule: when $\text{rank}(\hat{\ell}^{(h)}) \geq \text{rank}(\ell^{(h)})$, except for (IGNORE, IGNORE), memory influence is realized or excessive, so tokens with $d_{i,t}^{(h)} > 0$ serve as candidate locations. When $\text{rank}(\hat{\ell}^{(h)}) < \text{rank}(\ell^{(h)})$, actual use is insufficient. In this case, tokens with $d_{i,t}^{(h)} > 0$ may reflect the portion of memory influence already realized, so assigning the under-use penalty to them would suppress correct behavior. Instead, we focus on tokens with $d_{i,t}^{(h)} < 0$, whose likelihood the memory decreases but the model still generates, capturing cases in which content remains in the response despite a suppressive memory. Correct (IGNORE, IGNORE) has zero localization signal because it leaves no response footprint; its positive reward still contributes to sequence-level credit. Accordingly, we define

$$s^{(h)} = \begin{cases} 0, & (\ell^{(h)}, \hat{\ell}^{(h)}) = (\text{IGNORE}, \text{IGNORE}), \\ -1, & \text{rank}(\hat{\ell}^{(h)}) < \text{rank}(\ell^{(h)}), \\ +1, & \text{rank}(\hat{\ell}^{(h)}) \geq \text{rank}(\ell^{(h)}) \wedge (\ell^{(h)}, \hat{\ell}^{(h)}) \neq (\text{IGNORE}, \text{IGNORE}), \end{cases}$$

and $z_{i,t}^{(h)} = s^{(h)} d_{i,t}^{(h)}$, so that positive $z_{i,t}^{(h)}$ marks the t -th token as a candidate location for the credit of channel h . To make localization robust to extreme likelihood spikes and weak fluctuations, we first clip $z_{i,t}^{(h)}$ as $\bar{z}_{i,t}^{(h)} = \text{clip}(z_{i,t}^{(h)}, -d_{\max}, d_{\max})$. We then filter the localization signals at two levels: the signal-strength gate $G_i^{(h)} = \mathbf{1}[\max_t z_{i,t}^{(h)} > \delta_{\text{tok}}^{\text{abs}}]$ suppresses localization when no token-level signal exceeds the absolute threshold, while the channel-specific threshold $\delta^{(h)}$ filters out weak token-level fluctuations in the clipped scores, yielding

$$q_{i,t}^{(h)} = G_i^{(h)} \text{ReLU}(\bar{z}_{i,t}^{(h)} - \delta^{(h)}).$$

We maintain each channel-specific threshold $\delta^{(h)}$ using an exponential moving average (EMA) (see Appendix D.3 for details). The resulting $q_{i,t}^{(h)}$ captures a token-level manifestation of the atoms in channel h in the response and serves as a proxy signal for fine-grained advantage redistribution.

3.3. Channel-wise mean-preserving advantage redistribution

Having obtained token-level localization signals for each localizable channel, we normalize them within each response y_i of length T_i :

$$w_{i,t}^{(h)} = \frac{q_{i,t}^{(h)}}{Q_i^{(h)}}, \quad \text{where} \quad Q_i^{(h)} = \sum_{t=1}^{T_i} q_{i,t}^{(h)}.$$

The normalized weight $w_{i,t}^{(h)}$ determines the relative allocation of the credit of channel h across response tokens. To redistribute the channel advantage in these proportions while preserving the average advantage across the response, we assign a multiplier to each token. Requiring the multipliers to be proportional to $w_{i,t}^{(h)}$ and average to one across tokens yields $m_{i,t}^{(h)} = T_i w_{i,t}^{(h)}$ (see Appendix D.4 for the derivation and implementation details). For each channel h , the group-normalized advantage of response i is

$$A_i^{(h)} = \frac{R_i^{(h)} - \mu^{(h)}}{\sigma^{(h)} + \varepsilon},$$

where $\mu^{(h)}$ and $\sigma^{(h)}$ are the mean and standard deviation of the channel rewards across the rollout group. We then redistribute $A_i^{(h)}$ using the multiplier:

$$\tilde{A}_{i,t}^{(h)} = A_i^{(h)} \left[(1 - \eta_i^{(h)}) + \eta_i^{(h)} m_{i,t}^{(h)} \right].$$

Here $\eta_i^{(h)} = \chi_i^{(h)} \eta^{(h)}$, where $\eta^{(h)} \in [0, 1]$ is the channel-specific localization coefficient; setting $\eta^{(h)}$ to zero reduces the redistribution to uniform assignment of $A_i^{(h)}$ across response tokens. $\chi_i^{(h)}$ is the redistribution gate, defined as

$$\chi_i^{(h)} = \mathbf{1} \left[R_i^{(h)} A_i^{(h)} > 0 \wedge Q_i^{(h)} > 0 \right],$$

and permits token-level redistribution only when a nonzero localization signal is available and the channel reward agrees in sign with the channel advantage, preventing localized credit assignment from reversing the update direction specified by the channel reward. Finally, we aggregate and normalize the advantages across channels to maintain a stable learning scale and use the resulting token advantages in the clipped GRPO objective to update the policy (Appendix D.4).

4. Experiments

4.1. Setup

We compare MEMCALIB-RL with representative post-training methods, including SFT [21], two OPSD variants (OPSD-PG and OPSD-GKD) [39], GRPO [29], and GDPO [18]; we also report the original instruction-tuned models as Base. From the 13,500 training examples, we hold out a stratified subset of 1,500 as a validation set for hyperparameter selection. For SFT, we fine-tune the model using SFT data constructed from the remaining 12,000 examples (see Appendix E.1 for construction details). For OPSD, GRPO, GDPO, and MEMCALIB-RL, we initialize the model from the same cold-start checkpoint trained on SFT data constructed from a 4,000-example subset of the remaining 12,000 training examples and use the other 8,000 examples for subsequent training. To evaluate MEMCALIB-RL across model families and scales, we conduct experiments on Qwen3-8B [23], Ministral-3-8B-Instruct [17], and Qwen3.5-35B-A3B [24]. We evaluate all models on the held-out 1,500-example test set following Section 2.4, reporting the mean and standard deviation over three seeds. Appendix E.2 details the implementation of all methods.

4.2. Main results

As shown in Table 2, MEMCALIB-RL achieves the highest SCS and Exact across all three models, with improvements over the strongest baseline ranging from 1.73 to 7.29 on SCS and from 0.87 to 9.98 on Exact. The complete directional results in Table 9 show that, relative to Cold-start, all other post-training methods reduce errors in one direction while increasing errors in the other on at least one model. MEMCALIB-RL is the only method that reduces both over-use and under-use across all three models, demonstrating the effectiveness of using bidirectional counterfactual evidence to guide advantage redistribution across model families and scales. Moreover, without additional training, MEMCALIB-RL achieves the strongest overall performance on RPEval, showing that the gains transfer to an external memory-use benchmark (Appendix G).

4.3. Ablations and robustness

Having established the overall gains of MEMCALIB-RL, we next examine its localization design and the robustness of the resulting comparison across Judges. To this end, we conduct three complementary experiments on Qwen3-8B. First, to isolate the effects of localization granularity and direction-aware counterfactual evidence, we ablate both design choices. Sentence-level localization serves as a natural coarse-grained alternative: it segments the response into sentence units and applies one multiplier to all tokens within each unit, whereas token-level localization distinguishes individual tokens. At each granularity, we compare ordered bidirectional localization with an absolute-magnitude variant that retains only the magnitude of the counterfactual likelihood change (implementation details in Appendix H.2). As shown in Table 3, ordered bidirectional localization improves SCS and Exact by 1.59 and 1.80 points at the sentence level and by 4.84 and 6.76 points at the token level, with token-level ordered localization achieving the highest SCS and Exact scores. These results support our design in MEMCALIB-RL, which combines token-level localization with ordered bidirectional counterfactual evidence. Second, to assess sensitivity to localization strength, we evaluate $\eta^{(h)} \in \{0, 0.25, 0.5, 0.75, 1.0\}$ as the shared localization coefficient across the eight localizable channels, while holding all other settings fixed. Figure 3(a) shows that SCS and Exact both peak at $\eta^{(h)} = 0.75$, reaching 79.54 and 67.89, respectively. These results suggest that retaining a sequence-level component alongside token-level localization is beneficial. Further details are provided in Appendix H.1.

Finally, alongside the human evaluation of DeepSeek-V4-Pro in Appendix C, we re-evaluate the Qwen3-8B results in Table 2 using Qwen3.8-Max [27] as an alternative Judge. The results in Table 13 show that MEMCALIB-RL retains the highest SCS and Exact and the lowest sMUS and AUR among the eight methods. Method rankings also remain highly consistent with those obtained using DeepSeek-V4-Pro, with Spearman correlations of 0.976 for SCS and 1.000 for Exact (Table 14). These results show that MEMCALIB-RL’s advantage in overall memory use extends to evaluation by an alternative Judge. Appendix H.3 provides the full results and comparison procedure.

Table 2 | Main results across model families and scales (mean $\pm$ standard deviation over three seeds). Red and blue denote the best and second-best result for each model.

Method	Qwen3-8B		Ministral-3-8B-Instruct		Qwen3.5-35B-A3B	
	SCS $\uparrow$	Exact $\uparrow$	SCS $\uparrow$	Exact $\uparrow$	SCS $\uparrow$	Exact $\uparrow$
Base	31.17 $\pm$ 0.21	15.29 $\pm$ 0.23	31.88 $\pm$ 0.58	16.84 $\pm$ 0.48	26.54 $\pm$ 0.28	12.24 $\pm$ 0.25
Cold-start	51.98 $\pm$ 1.24	33.87 $\pm$ 2.19	66.87 $\pm$ 1.05	50.62 $\pm$ 1.56	71.52 $\pm$ 0.96	56.56 $\pm$ 1.20
SFT	65.72 $\pm$ 0.64	49.49 $\pm$ 1.24	70.94 $\pm$ 0.28	55.89 $\pm$ 0.52	75.68 $\pm$ 0.46	61.87 $\pm$ 0.71
OPSD-PG	54.36 $\pm$ 0.60	35.80 $\pm$ 0.60	70.48 $\pm$ 0.42	54.71 $\pm$ 0.30	71.29 $\pm$ 0.25	55.64 $\pm$ 0.87
OPSD-GKD	57.39 $\pm$ 0.41	39.49 $\pm$ 0.38	71.34 $\pm$ 0.18	55.91 $\pm$ 0.31	71.67 $\pm$ 0.85	56.89 $\pm$ 1.55
GRPO	67.61 $\pm$ 0.71	52.18 $\pm$ 0.82	77.21 $\pm$ 0.46	65.44 $\pm$ 0.80	78.25 $\pm$ 0.87	66.60 $\pm$ 1.28
GDPO	72.25 $\pm$ 0.30	57.91 $\pm$ 0.40	76.98 $\pm$ 0.35	63.84 $\pm$ 0.43	79.39 $\pm$ 1.02	67.51 $\pm$ 1.44
MEMCALIB-RL	79.54 $\pm$ 0.77	67.89 $\pm$ 1.07	79.00 $\pm$ 0.19	66.31 $\pm$ 0.60	81.12 $\pm$ 0.78	70.16 $\pm$ 1.23

Table 3 | Localization ablations on Qwen3-8B (mean $\pm$ standard deviation over three seeds). Red and blue denote the best and second-best result in each column.

Granularity	Signal	SCS $\uparrow$	Exact $\uparrow$	sMOS $\downarrow$	sMUS $\downarrow$	AOR $\downarrow$	AUR $\downarrow$
Sentence	Absolute	76.59 $\pm$ 0.11	63.87 $\pm$ 0.27	13.70 $\pm$ 0.45	11.65 $\pm$ 0.27	21.53 $\pm$ 0.75	19.80 $\pm$ 0.53
Sentence	Ordered	78.18 $\pm$ 0.04	65.67 $\pm$ 0.53	12.98 $\pm$ 0.19	10.52 $\pm$ 0.18	20.60 $\pm$ 0.55	18.20 $\pm$ 0.44
Token	Absolute	74.70 $\pm$ 0.78	61.13 $\pm$ 0.98	14.40 $\pm$ 0.39	12.96 $\pm$ 0.48	22.20 $\pm$ 0.55	22.11 $\pm$ 0.68
Token	Ordered	79.54 $\pm$ 0.77	67.89 $\pm$ 1.07	14.98 $\pm$ 0.60	6.81 $\pm$ 0.28	23.49 $\pm$ 0.74	12.29 $\pm$ 0.57

4.4. Mechanism analysis

We then assess whether counterfactual localization directs credit to memory-influenced tokens and analyze the training dynamics of MEMCALIB-RL. For localization analysis, we sample 20 response-channel pairs from each of the eight localizable channels. For each sampled pair, we first aggregate token-level credit within each sentence, and a human annotator then labels each sentence as IRRELEVANT, PARTIAL, or KEY. Using these annotations, we compute, for each pair, the fraction of response tokens contained in sentences of each label and the fraction of total credit assigned to those sentences, and record a Top-3 localization hit if at least one of the three highest-credit sentences is labeled PARTIAL or KEY. Across all eight channels, Figure 3(h) shows that IRRELEVANT, PARTIAL, and KEY sentences account for 48.0%, 17.4%, and 34.6% of response tokens but receive 24.7%, 13.3%, and 61.9% of the total credit, respectively. The Top-3 hit rate is 99.0% for the five support-targeting channels (AB^- , AC^- , B^+ , BC^- , and C^+) and 96.7% for the three under-use channels (BA^- , CA^- , and CB^-). These results show that counterfactual localization indeed assigns more credit to memory-influenced content in the response. Annotation and aggregation details, together with representative cases from both channel groups, are provided in Appendix I.

To analyze training dynamics, we divide the eight localizable channels into a correct-use group (B^+ and C^+) and a corrective group (the other six). Figure 3(b) shows that the B^+ and C^+ trigger rates rise, whereas five of the six corrective channels decline, with the low-frequency BC^- channel remaining stable. This trend indicates that the policy progressively shifts from memory-use errors toward correct use. Figure 3(c) shows that the token-level localization rate falls from 37.9% to 16.4% for correct-use channels but only from 87.7% to 83.3% for corrective channels. Figure 3(d) explains the declining localization rates: the non-degenerate-group rates of both channel groups decrease, showing that the rollouts for a query increasingly receive identical channel rewards. As a result, the redistribution gate more frequently falls back to sequence-level credit. Although policy entropy initially falls in parallel, it recovers late while the non-degenerate-group rates remain low, showing that the loss of within-group reward variation reflects convergence in memory-use outcomes rather than entropy collapse. Figure 3(e) further shows that group degeneration mainly reflects uniformly correct memory use in correct-use channels and the uniform absence of the corresponding error in corrective




Figure 3 | Sensitivity and mechanism analyses on Qwen3-8B. (a) Sensitivity to the localization coefficient η ; (b–g) training dynamics of MEMCALIB-RL; and (h) token and credit shares across human-annotated sentence roles for all eight localizable channels. Curves in (b), (d), (f), and (g) use a seven-step centered mean; error bars in (a) show standard deviations over three seeds and those in (h) show 95% bootstrap confidence intervals.

channels. Together, Figures 3(b–e) show that training yields more appropriate and consistent memory use, while the mechanism correctly falls back when comparative signals vanish. Finally, Figures 3(f–g) show that channel thresholds separate into distinct empirical scales, supporting channel-adaptive filtering, while the token-weighted mean sentence multiplier remains 1 as its distribution widens, confirming differentiated local credit without response-level scale drift. Full analysis details and supporting statistics for Figure 3(b–g) are provided in Appendix J.

5. Conclusion and Limitations

Agent memory is effective only if the LLM gives each supplied proposition the appropriate level of influence. Our results show that this capability cannot be taken for granted: frontier models frequently mismatch target and actual use levels, while common post-training algorithms can improve one error direction at the expense of the other. MEMCALIB makes this overlooked problem measurable, while MEMCALIB-RL addresses it through bidirectional counterfactual credit localization and advantage redistribution, yielding stronger overall memory-use performance and a better balance between over-use and under-use.

Counterfactual localization provides a proxy for locating channel-level credit within a response, but its token-level attribution may not always be accurate. However, the algorithm’s robust design mitigates this limitation through sequence-level fallback and channel-specific control of localization strength via $\eta^{(h)}$; experiments across model families and scales, together with human evaluation, demonstrate the practical effectiveness of counterfactual localization. Developing more principled and accurate ways to use this signal remains an important direction for future work.

Acknowledgments

This work was supported by Qwen Business Unit through Alibaba Research Intern Program.

References

- [1] Z. Cao, J. Deng, L. Yu, W. Zhou, Z. Liu, B. Ding, and H. Zhao. Remember me, refine me: A dynamic procedural memory framework for experience-driven agent evolution. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 16803–16822, 2026.

-
- [2] DeepSeek-AI. Deepseek-v4: Towards highly efficient million-token context intelligence. *arXiv preprint arXiv:2606.19348*, 2026.
 - [3] N. Ding, Y. Chen, B. Xu, Y. Qin, S. Hu, Z. Liu, M. Sun, and B. Zhou. Enhancing chat language models by scaling high-quality instructional conversations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 3029–3051, 2023. doi: 10.18653/v1/2023.emnlp-main.183.
 - [4] X. Feng, W. Gan, X. Chen, Q. Dai, and Y. Liu. How does personalized memory shape LLM behavior? benchmarking rational preference utilization in personalized assistants. *arXiv preprint arXiv:2601.16621*, 2026.
 - [5] X. He, S. Chen, Z. Ju, X. Dong, H. Fang, S. Wang, Y. Yang, J. Zeng, R. Zhang, R. Zhang, M. Zhou, P. Zhu, and P. Xie. MedDialog: Two large-scale medical dialogue datasets. *arXiv preprint arXiv:2004.03329*, 2020. URL <https://arxiv.org/abs/2004.03329v2>.
 - [6] D. Hendrycks, S. Basart, S. Kadavath, M. Mazeika, A. Arora, E. Guo, C. Burns, S. Puranik, H. He, D. Song, and J. Steinhardt. Measuring coding challenge competence with APPS. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, volume 1, 2021. URL <https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/c24cd76e1ce41366a4bbe8a49b02a028-Abstract-round2.html>.
 - [7] M. Honnibal, I. Montani, S. Van Landeghem, and A. Boyd. spaCy: Industrial-strength natural language processing in Python, 2020. URL <https://doi.org/10.5281/zenodo.1212303>.
 - [8] Y. Hu, S. Liu, Y. Yue, G. Zhang, B. Liu, F. Zhu, J. Lin, H. Guo, S. Dou, Z. Xi, et al. Memory in the age of AI agents. *arXiv preprint arXiv:2512.13564*, 2025.
 - [9] Y. Hu, Z. Long, J. Guo, X. Sui, X. Fu, W. Zhao, Y. Zhao, and B. Qin. OP-Bench: Benchmarking over-personalization for memory-augmented personalized conversational agents. *arXiv preprint arXiv:2601.13722*, 2026.
 - [10] Y. Hu, Y. Wang, and J. McAuley. Evaluating memory in LLM agents via incremental multi-turn interactions. In *International Conference on Learning Representations*, 2026.
 - [11] S. Kim, J. Suk, S. Longpre, B. Y. Lin, J. Shin, S. Welleck, G. Neubig, M. Lee, K. Lee, and M. Seo. Prometheus 2: An open source language model specialized in evaluating other language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 4334–4353, 2024. doi: 10.18653/v1/2024.emnlp-main.248.
 - [12] A. Köpf, Y. Kilcher, D. von Rütte, S. Anagnostidis, Z. R. Tam, K. Stevens, A. Barhoum, D. Nguyen, O. Stanley, R. Nagyfi, S. ES, S. Suri, D. Glushkov, A. Dantuluri, A. Maguire, C. Schuhmann, H. Nguyen, and A. Mattick. OpenAssistant conversations - democratizing large language model alignment. In *Advances in Neural Information Processing Systems*, 2023.
 - [13] N. Lambert, L. Tunstall, N. Rajani, and T. Thrush. Hugging face H4 stack exchange preference dataset, 2023. URL <https://huggingface.co/datasets/HuggingFaceH4/stack-exchange-preferences>.
 - [14] D. Li, B. Jiang, L. Huang, A. Beigi, C. Zhao, Z. Tan, A. Bhattacharjee, Y. Jiang, C. Chen, T. Wu, K. Shu, L. Cheng, and H. Liu. From generation to judgment: Opportunities and challenges of LLM-as-a-judge. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 2757–2791, 2025. doi: 10.18653/v1/2025.emnlp-main.138.
 - [15] K. Li, X. Yu, Z. Ni, Y. Zeng, Y. Xu, Z. Zhang, X. Li, J. Sang, X. Duan, X. Wang, C. Liu, and J. Tan. TiMem: Temporal-hierarchical memory consolidation for long-horizon conversational agents. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 21700–21720, 2026. URL <https://dblp.org/rec/conf/acl/LiYNZXZLSDWLT26>.
 - [16] Y. Li, Z. Li, K. Zhang, R. Dan, S. Jiang, and Y. Zhang. ChatDoctor: A medical chat model fine-tuned on a large language model meta-AI (LLaMA) using medical domain knowledge. *arXiv preprint arXiv:2303.14070*, 2023. URL <https://arxiv.org/abs/2303.14070v5>.
-

- [17] A. H. Liu, K. Khandelwal, S. Subramanian, V. Jouault, A. Rastogi, A. Sadé, A. Jeffares, A. Jiang, A. Cahill, A. Gavaudan, et al. Ministral 3. *arXiv preprint arXiv:2601.08584*, 2026.
- [18] S.-Y. Liu, X. Dong, X. Lu, S. Diao, P. Belcak, M. Liu, M.-H. Chen, H. Yin, Y.-C. F. Wang, K.-T. Cheng, Y. Choi, J. Kautz, and P. Molchanov. GDPO: Group reward-decoupled normalization policy optimization for multi-reward RL optimization. *arXiv preprint arXiv:2601.05242*, 2026.
- [19] J. Luo, Y. Tian, C. Cao, Z. Luo, H. Lin, K. Li, C. Kong, R. Yang, and J. Ma. From storage to experience: A survey on the evolution of LLM agent memory mechanisms. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 41622–41652, 2026.
- [20] A. Maharana, D.-H. Lee, S. Tulyakov, M. Bansal, F. Barbieri, and Y. Fang. Evaluating very long-term conversational memory of LLM agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 13851–13870, 2024. doi: 10.18653/v1/2024.acl-long.747.
- [21] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. L. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, J. Schulman, J. Hilton, F. Kelton, L. Miller, M. Simens, A. Askell, P. Welinder, P. F. Christiano, J. Leike, and R. Lowe. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744, 2022.
- [22] J. S. Park, J. C. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, pages 2:1–2:22, 2023. doi: 10.1145/3586183.3606763.
- [23] Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [24] Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026. URL <https://qwen.ai/blog?id=qwen3.5>.
- [25] Qwen Team. Qwen3.7: The agent frontier, May 2026. URL <https://qwen.ai/blog?id=qwen3.7>.
- [26] Qwen Team. Qwen3.7-Plus: Multimodal agent intelligence, May 2026. URL <https://qwen.ai/blog?id=qwen3.7-plus>.
- [27] Qwen Team. Qwen3.8-Max: A new bar for coding and cowork, August 2026. URL <https://qwen.ai/blog?id=qwen3.8>.
- [28] C. Sewell. markdown-it-py: Python port of markdown-it, 2025. URL <https://github.com/executablebooks/markdown-it-py/releases/tag/v4.0.0>. Version 4.0.0.
- [29] Z. Shao, P. Wang, Q. Zhu, R. Xu, J. Song, X. Bi, H. Zhang, M. Zhang, Y. K. Li, Y. Wu, and D. Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. URL <https://arxiv.org/abs/2402.03300v3>.
- [30] Y. Shen, K. Li, W. Zhou, and S. Hu. Mem2ActBench: A benchmark for evaluating long-term memory utilization in task-oriented autonomous agents. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pages 8173–8190, 2026.
- [31] Y. Wei, Z. Wang, J. Liu, Y. Ding, and L. Zhang. Magicoder: Empowering code generation with OSS-Instruct. In *Proceedings of the 41st International Conference on Machine Learning*, pages 52632–52657, 2024. URL <https://dblp.org/rec/conf/icml/0003W0D024>.
- [32] D. Wu, H. Wang, W. Yu, Y. Zhang, K.-W. Chang, and D. Yu. LongMemEval: Benchmarking chat assistants on long-term interactive memory. In *International Conference on Learning Representations*, 2025. URL <https://dblp.org/rec/conf/iclr/WuWYZCY25>.
- [33] Y. Wu, T. Wu, M. Zhu, H. Sha, and H. Wang. StratMem-Bench: Evaluating strategic memory use in virtual character conversation beyond factual recall. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pages 32309–32328, 2026.

- [34] Z. Xiang, Z. Chen, Y. Tang, Z. Wei, R. Ning, Y. Lin, Q. Zhang, and J. Su. MemSyco-Bench: Benchmarking sycophancy in agent memory. *arXiv preprint arXiv:2607.01071*, 2026.
- [35] Z. Xiong, Y. Lin, W. Xie, P. He, Z. Liu, J. Tang, H. Lakkaraju, and Z. Xiang. How memory management impacts LLM agents: An empirical study of experience-following behavior. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pages 623–645, 2026.
- [36] S. Yan, X. Yang, Z. Huang, E. Nie, Z. Ding, Z. Li, X. Ma, J. Bi, K. Kersting, J. Z. Pan, H. Schuetze, V. Tresp, and Y. Ma. Memory-R1: Enhancing large language model agents to manage and utilize memories via reinforcement learning. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pages 12805–12825, 2026.
- [37] A. Yehudai, L. Eden, A. Li, G. Uziel, Y. Zhao, R. Bar-Haim, A. Cohan, and M. Shmueli-Scheuer. A survey on evaluation of LLM-based agents. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 26690–26714, 2026. doi: 10.18653/v1/2026.findings-acl.1330.
- [38] S. Yoon, S. Kim, H. Hong, W. Jeung, Y. Kim, W. Seo, H. Yeen, and A. No. BenchPreS: A benchmark for context-aware personalized preference selectivity of persistent-memory LLMs. *arXiv preprint arXiv:2603.16557*, 2026.
- [39] S. Zhao, Z. Xie, M. Liu, J. Huang, G. Pang, F. Chen, and A. Grover. Self-distilled reasoner: On-policy self-distillation for large language models. *arXiv preprint arXiv:2601.18734*, 2026.
- [40] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, and I. Stoica. Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information Processing Systems*, 2023. URL <https://dblp.org/rec/conf/nips/ZhengC00WZL0LXZ23>.
- [41] W. Zhong, L. Guo, Q. Gao, H. Ye, and Y. Wang. MemoryBank: Enhancing large language models with long-term memory. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 19724–19731, 2024. doi: 10.1609/aaai.v38i17.29946.

A. Related Work

Agent memory and evaluation. Agent memory research focuses on the formation, updating, and retrieval of factual, experiential, and working memory [8, 19]. LongMemEval and LoCoMo evaluate long-term conversational memory and reasoning [32, 20]. MemoryAgentBench assesses retrieval, test-time learning, long-range understanding, and selective forgetting through incremental interactions [10]; Mem2ActBench evaluates memory-driven tool selection and parameter grounding [30]. Memory-R1 uses reinforcement learning to optimize memory management and memory-based answer generation [36], while ReMe distills, reuses, and refines procedural experience [1].

Selective and personalized memory use. OP-Bench and MemSyc-Bench expose over-personalization and memory-induced sycophancy [9, 34]. RPEval assigns user preferences IGNORE, SUPPORT, or DOMINATE strategies and evaluates personalization errors arising from mismatches between intended and actual use [4]. BenchPreS evaluates whether stored preferences are appropriately applied or suppressed in third-party communication contexts [38]. StratMem-Bench distinguishes required, supportive, and irrelevant memories in virtual-character dialogue and assesses item-level use [33]. MEMCALIB evaluates whether models use memory appropriately across health, general assistance, and coding, with atomic propositions embedded in natural composite memory blocks. Atom-specific rubrics assess each proposition’s actual influence on the response against its target level, yielding measures of over-use, under-use, and overall memory-use performance. Table 4 summarizes the evaluation scope and granularity of these benchmarks.

Table 4 | Comparison of memory-use benchmarks by evaluation scope and granularity. ✓ indicates an explicit evaluation task or criterion; × indicates that the aspect is not explicitly evaluated. Conflict handling covers conflicts with task requirements, facts, or other memories. Item-level evaluation assesses individual memories or preferences.

Benchmark	Setting	Applicability	Conflict handling	Per-item use	Memory-use evaluation unit
OP-Bench [9]	Personalized dialogue	✓	✓	×	Response; response set
MemSyc-Bench [34]	Personalized QA and decisions	✓	✓	×	Response
RPEval [4]	Everyday assistance	✓	✓	✓	Individual preference
BenchPreS [38]	Third-party communication	✓	✓	✓	Preference attribute
StratMem-Bench [33]	Virtual-character dialogue	✓	×	✓	Individual memory item
MEMCALIB (ours)	Health, general assistance, coding	✓	✓	✓	Atomic proposition within a memory block

B. MEMCALIB Construction and Data Statistics

B.1. Construction and quality control

We construct MEMCALIB from eight public question-answering datasets spanning health dialogue, general assistance, and coding. Table 5 lists the source datasets and summarizes the tasks they cover in each domain. We derive each sample using a six-stage pipeline that is iteratively refined based on human review of fresh pilot batches and then frozen for full-scale construction. Figure 4 illustrates both the six-stage pipeline and its refinement through human feedback. During data construction, we use Qwen3.7-Max [25] with stage-specific prompts for generation and annotation.

Table 5 | Public datasets used to construct MEMCALIB.

Domain	Source datasets	Coverage
Health dialogue	MedDialog [5]; ChatDoctor–HealthCareMagic [16]	Symptoms, medication, chronic conditions, and care seeking
General assistance	UltraChat 200K [3]; OpenAssistant OASST1 and OASST2 [12]	Explanation, planning, writing, and everyday assistance
Coding	Stack Exchange Preferences [13]; APPS [6]; Magicoder OSS-Instruct [31]	Algorithms, APIs and libraries, systems, and debugging

1. Source screening. We first assess source samples for query–memory derivation through deterministic filtering followed by LLM-based semantic screening. Deterministic filtering checks that sample identifiers, source information, questions, and reference answers are present, enforces English-language and length requirements, and removes noisy text and exact or near-duplicate samples. Semantic screening checks whether the source question and reference answer form a coherent, informative QA pair, and whether the question or context contains explicit information suitable for memory extraction.

2. Query and memory derivation. For each admitted source sample, we prompt the LLM to extract facts, preferences, constraints, and past events explicitly stated in the question or context, and group related details into memory passages for subsequent atomic annotation. The LLM rewrites the question to preserve the original request while removing the information assigned to these passages. The reference answer aids task understanding during this process. During generation, we instruct the LLM to produce a query that neither states nor implies any extracted memory detail and remains coherent and answerable with the extracted memories. We verify both conditions during the subsequent independent LLM-based review described in Stage 6.

3. Atomic annotation. We next use the LLM to decompose the extracted memory passages into independently judgeable atomic propositions and assign each atom an ideal use level: no answer-specific footprint (IGNORE), bounded local support (BOUND), or control over a material conclusion, constraint, or recommendation (CONTROL). To ensure that each sample tests the use of task-relevant memory, we retain only samples containing at least one BOUND or CONTROL atom.

4. Controlled distractor construction. Because source questions and their context primarily contain information relevant to the original request, the LLM assigns most extracted atoms a BOUND or CONTROL level in Stage 3. To assess models’ ability to ignore distracting memories, we prompt the LLM to generate one hard IGNORE atom and a variable number of additional IGNORE atoms for each retained item. The LLM constructs the hard atom to appear relevant through topical proximity or same-user plausibility while keeping its content outside the current answer scope. It generates the additional atoms to vary memory-block length and composition and simulate long-tail retrieval noise. After adding these distractors, we prompt the LLM to generate an atom-specific rubric for every source-derived and added atom. Each rubric specifies observable response-text criteria for distinguishing the three actual use levels.

5. Block composition and naturalization. After adding the distractors and generating a rubric for each atom, we assemble the atoms into memory blocks using predefined counts for the blocks in each sample and the atoms within each block. These counts vary across samples, with atom counts also varying across blocks within each sample, to cover different memory-context sizes and block complexities. We preserve atom order and allow atoms with different ideal use levels to coexist within a block, to reflect how a retrieved memory can mix useful context with irrelevant details. We then prompt the LLM to rewrite each block as a coherent natural-language paragraph, preserving every proposition’s meaning without omitting atoms or introducing additional information.




Figure 4 | Overview of the MEMCALIB construction pipeline and its refinement through human feedback. During pipeline development, each iteration constructs a fresh batch of 50 pilot samples balanced across the three domains. Human review identifies sample-level issues and guides revisions to the relevant stages and prompts. This process continues until the review pass rate reaches at least 90%, after which the pipeline is frozen for full-scale construction. Solid arrows indicate the sequence of data construction steps, while dashed arrows indicate human feedback for pipeline refinement.

6. Quality assurance. After naturalizing the memory blocks, we validate each completed sample. We first run deterministic checks to verify that each sample contains the query, memory blocks, and atomic annotations, that each atom has a valid ideal use level and a complete rubric, and that every atom belongs to exactly one block in its recorded order. We then prompt Qwen3.7-Max and DeepSeek-V4-Pro [2] to review each sample independently. Each LLM checks whether source-derived atoms faithfully capture the original context and express distinct propositions; whether the query excludes extracted memory details and remains coherent and answerable with the memories; whether rewritten blocks preserve every proposition without adding information; and whether ideal levels reflect each atom’s role in the task and rubrics provide observable response-text criteria for distinguishing actual use levels. We accept a sample through semantic review only when both LLMs approve it.

Iterative refinement through human feedback. Before full-scale construction, we iteratively refine the six-stage pipeline through human review. At each iteration, we construct a fresh batch of 50 pilot samples balanced across health, general assistance, and coding. We manually inspect all 50 pilot samples, including their query-memory splits, ideal use levels, and rubric criteria. A sample passes only if no issues are identified across all reviewed aspects. For samples that fail, we analyze the identified issues, trace them to the relevant construction stages, and revise the corresponding procedures and LLM prompts. We repeat this process with a newly constructed batch at each iteration until the sample-level human-review pass rate reaches at least 90%. We then freeze the pipeline and use it for full-scale benchmark construction.

The [example below](#) illustrates a constructed health-domain sample, including its query, memory blocks, ideal use levels, and rubric criteria for assessing each atom’s influence on the response. For space, we show only selected rubric fields for one atom at each ideal use level.

An example of MEMCALIB

Current query x

What medical evaluations are typically recommended for a couple who has been having frequent unprotected intercourse without conceiving?

Memory blocks M

b_1 The user is 21 years old and has been attempting to conceive through daily unprotected intercourse for three years. The user’s partner is 24 years old, and the couple resides in a suburban apartment complex.

b_2 The user reports irregular menstrual cycles and tracks menstrual cycles using a paper calendar.

b_3 The user is currently writing a master’s thesis on the sociological impacts of infertility treatment access.

b_4 The user acquired an indigo coin display tray and stored it in the foyer sideboard drawer in May 2008, having lined the tray with acid-free felt before placing it in that drawer in May 2008.

Evaluation annotations (not provided to the model)

Atom	Atomic proposition	Ideal level
$m_{1,1}$	The user is 21 years old.	BOUND
$m_{1,2}$	The user has been attempting to conceive through daily unprotected intercourse for three years.	CONTROL
$m_{1,3}$	The user’s partner is 24 years old.	IGNORE
$m_{1,4}$	The couple resides in a suburban apartment complex.	IGNORE
$m_{2,1}$	The user reports irregular menstrual cycles.	CONTROL
$m_{2,2}$	The user tracks menstrual cycles using a paper calendar.	IGNORE
$m_{3,1}$	The user is currently writing a master’s thesis on the sociological impacts of infertility treatment access.	IGNORE
$m_{4,1}$	The user acquired an indigo coin display tray and stored it in the foyer sideboard drawer in May 2008.	IGNORE
$m_{4,2}$	The indigo coin display tray was lined with acid-free felt before being placed in the foyer sideboard drawer in May 2008.	IGNORE

Rubric examples $m_{3,1}$ — **Ideal level: IGNORE**

Expected behavior. Answer provides a direct, clinically focused list of evaluations without referencing the user’s academic work or sociological perspectives.

Over-use. Inclusion of academic framing, sociological terminology, or references to the thesis in the medical advice.

 $m_{1,1}$ — **Ideal level: BOUND**

Expected behavior. Answer references user’s young age to contextualize evaluation recommendations or prognosis without altering core test list.

Under-use. No acknowledgment of age despite its relevance to fertility counseling.

Over-use. Attributing infertility solely to age or recommending different tests based only on age.

 $m_{1,2}$ — **Ideal level: CONTROL**

Expected behavior. Answer explicitly recommends medical fertility evaluation for both partners based on duration of unprotected intercourse.

Under-use. Failing to recommend evaluation despite three-year history.

B.2. Benchmark composition

MEMCALIB contains 15,000 examples across health, general assistance, and coding. Table 6 reports the total numbers of examples, memory blocks, and atomic propositions in each domain, while Table 7 details the distribution of ideal use levels and the structure of the memory blocks. A mixed-target block contains atoms with at least two different ideal use levels. The percentages indicate the share of atoms at each ideal level and the share of examples containing a mixed-target block. The ranges give the minimum and maximum numbers of atoms per block and per example. To support consistent response-text evaluation across domains, we formulate coding queries as natural-language tasks covering implementation planning, behavior prediction, code understanding, and debugging diagnosis, allowing us to apply atom-specific rubrics directly to the generated responses.

Table 6 | MEMCALIB composition by domain. Block and proposition counts are totals over examples in each domain.

Domain	Examples	Memory blocks	Atomic propositions
Health	7,500	37,400	117,110
General assistance	3,750	18,700	58,495
Coding	3,750	18,700	58,615
Total	15,000	74,800	234,220

Table 7 | Ideal use levels and memory structure in MEMCALIB.

Statistic	Value
Atoms with ideal level IGNORE	197,510 (84.3%)
Atoms with ideal level BOUND	18,943 (8.1%)
Atoms with ideal level CONTROL	17,767 (7.6%)
Mixed-target blocks	20,097
Examples containing a mixed-target block	14,949 (99.66%)
Atoms per block (min-max)	1-20
Atoms per example (min-max)	6-63

C. Human Evaluation of the Judge

To assess how reliably DeepSeek-V4-Pro judges the actual use level of each memory atom in a model response, we sample 300 response-atom pairs for human evaluation. Of these, 150 are randomly sampled to measure agreement under the natural evaluation distribution. The other 150 are stratified across the nine combinations of ideal and Judge-assessed actual use levels to examine less frequent combinations and use-level boundaries. For each pair, a human annotator applies the same atom-specific rubric, without seeing the Judge’s decision, to assess its actual use level as IGNORE, BOUND, or CONTROL. The results in Figure 5(a) show strong agreement under the natural evaluation distribution: the Judge and human annotator assess the same actual use level in 145 of 150 cases, yielding 96.7% exact agreement (95% Wilson CI: 92.4%–98.6%), with Cohen’s $\kappa = 0.872$. All 127 human-assessed IGNORE cases match the Judge’s assessment. Among the five disagreements, three lie at the BOUND/CONTROL boundary; the Judge assesses a lower actual use level in four cases and a higher one in one. The stress-stratified results in Figure 5(b) yield 111/150 exact agreement (74.0%; 95% Wilson CI: 66.4%–80.4%), with Cohen’s $\kappa = 0.610$. The largest concentration of disagreements again lies at the BOUND/CONTROL boundary: 16 cases are assessed as BOUND by the human annotator and CONTROL by the Judge, with 10 in the reverse direction. Together, these results support the reliability of DeepSeek-V4-Pro as the Judge under the natural evaluation distribution, while identifying the distinction between bounded and controlling influence as the primary remaining uncertainty.




Figure 5 | Confusion matrices comparing human and DeepSeek-V4-Pro assessments of actual use levels. Rows correspond to human assessments and columns to the Judge’s assessments; cells show counts and row-normalized percentages. (a) Natural-distribution subset. (b) Stress-stratified subset.

D. MEMCALIB-RL Details

D.1. Reward channels and localization targets

Table 8 summarizes the notation, per-atom rewards, and localization targets for all nine ideal–actual use transitions. The superscripts $+$ and $-$ in the channel labels indicate the sign of the reward. The sign of d indicates whether the channel’s memory atoms increase or decrease the likelihood of a generated token relative to removing those atoms.

Table 8 | Nine ideal–actual reward channels, shorthand notation, per-atom rewards, and localization targets.

Notation	Ideal ($\ell^{(h)}$) $\rightarrow$ actual ($\hat{\ell}^{(h)}$)	Per-atom reward	Localization target
A^+	IGNORE $\rightarrow$ IGNORE	+1	sequence-level
AB^-	IGNORE $\rightarrow$ BOUND	-1	supported tokens ($d > 0$) or fallback
AC^-	IGNORE $\rightarrow$ CONTROL	-2	supported tokens ($d > 0$) or fallback
BA^-	BOUND $\rightarrow$ IGNORE	-1	suppressed tokens ($d < 0$) or fallback
B^+	BOUND $\rightarrow$ BOUND	+1	supported tokens ($d > 0$) or fallback
BC^-	BOUND $\rightarrow$ CONTROL	-1	supported tokens ($d > 0$) or fallback
CA^-	CONTROL $\rightarrow$ IGNORE	-2	suppressed tokens ($d < 0$) or fallback
CB^-	CONTROL $\rightarrow$ BOUND	-1	suppressed tokens ($d < 0$) or fallback
C^+	CONTROL $\rightarrow$ CONTROL	+1	supported tokens ($d > 0$) or fallback

D.2. Group normalization and policy objective

GRPO aggregates channel rewards before group normalization, whereas GDPO normalizes each channel separately before aggregation. For a rollout group of G responses to the same query–memory pair (x, M) , GRPO computes the total reward $R_i = \sum_h R_i^{(h)}$ for the i -th response and normalizes it within the group [29]:

$$A_i^{\text{GRPO}} = \frac{R_i - \mu_R}{\sigma_R + \epsilon},$$

where μ_R and σ_R are the mean and standard deviation of the total rewards across the G responses:

$$\mu_R = \frac{1}{G} \sum_{i=1}^G R_i, \quad \sigma_R = \sqrt{\frac{1}{G} \sum_{i=1}^G (R_i - \mu_R)^2}.$$

Here $\varepsilon > 0$ is a small constant included in the denominators to avoid division by zero. GRPO then assigns this response-level advantage uniformly to all tokens in y_i :

$$A_{i,t}^{\text{GRPO}} = A_i^{\text{GRPO}}, \quad t = 1, \dots, T_i,$$

where T_i is the length of response y_i .

GDPO first normalizes each channel's rewards across the G responses in a rollout group to obtain a channel-specific advantage [18]:

$$A_i^{(h)} = \frac{R_i^{(h)} - \mu^{(h)}}{\sigma^{(h)} + \varepsilon},$$

where $\mu^{(h)}$ and $\sigma^{(h)}$ are the mean and standard deviation of the rewards of channel h within that group:

$$\mu^{(h)} = \frac{1}{G} \sum_{i=1}^G R_i^{(h)}, \quad \sigma^{(h)} = \sqrt{\frac{1}{G} \sum_{i=1}^G (R_i^{(h)} - \mu^{(h)})^2}.$$

After computing these advantages within each rollout group, GDPO sums them across channels for each response in the update batch:

$$S_j = \sum_h A_j^{(h)}, \quad j \in \mathcal{B},$$

where $\mathcal{B}$ indexes all valid responses across the rollout groups in the update batch. For the j -th batch response, $A_j^{(h)}$ denotes the channel advantage computed within its own rollout group. These summed advantages can grow in magnitude as more channels are combined, so the GDPO baseline applies a second normalization across all responses in $\mathcal{B}$ using root-mean-square (RMS) scaling to maintain a stable scale:

$$A_j^{\text{GDPO}} = \frac{S_j}{\sigma_S}, \quad (1)$$

where σ_S is the regularized RMS of the summed advantages over the update batch:

$$\sigma_S = \sqrt{\frac{1}{|\mathcal{B}|} \sum_{j \in \mathcal{B}} S_j^2 + \varepsilon}.$$

We set $\varepsilon = 10^{-6}$ to prevent division by zero. GDPO also assigns this response-level advantage uniformly to all tokens in y_j :

$$A_{j,t}^{\text{GDPO}} = A_j^{\text{GDPO}}, \quad t = 1, \dots, T_j,$$

where T_j is the length of response y_j .

Both methods use the resulting token advantages in the same clipped policy objective. For either method, let $A_{j,t}$ denote the token advantage for the j -th response y_j in the update batch, θ_{old} the frozen rollout policy, and θ the policy being updated. The importance ratio is

$$\rho_{j,t}(\theta) = \frac{\pi_{\theta}(y_{j,t} \mid x, M, y_{j,<t})}{\pi_{\theta_{\text{old}}}(y_{j,t} \mid x, M, y_{j,<t})}.$$

Here (x, M) is the query–memory pair associated with response y_j . The policy objective averages the clipped terms over all response tokens in the update batch:

$$J(\theta) = \frac{1}{\sum_{j \in \mathcal{B}} T_j} \sum_{j \in \mathcal{B}} \sum_{t=1}^{T_j} \min(\rho_{j,t}(\theta) A_{j,t}, \text{clip}(\rho_{j,t}(\theta), 1 - \epsilon_{\text{clip}}, 1 + \epsilon_{\text{clip}}) A_{j,t}) - \beta_{\text{KL}} J_{\text{KL}}(\theta). \quad (2)$$

Here $\epsilon_{\text{clip}} > 0$ is the clipping radius. $J_{\text{KL}}(\theta)$ is the KL penalty to the reference policy and β_{KL} is its coefficient. The policy parameters are updated by maximizing $J(\theta)$.

D.3. Channel-specific threshold initialization and updates

We maintain a separate token threshold for each of the eight localizable channels in Section 3.2. For the r -th rollout batch, let $\delta_r^{(h)}$ denote the threshold for channel h . We initialize every channel's threshold to $\delta_1^{(h)} = 0.02$. Within the r -th batch, all responses use the same threshold $\delta_r^{(h)}$ for channel h . We update it after processing the batch.

To update the threshold for channel h , we pool all direction-aligned, clipped token scores $\bar{z}_{j,t}^{(h)}$ from successfully scored counterfactuals in the r -th batch into $\mathcal{P}_r^{(h)}$. When this pool contains at least 256 scores, we compute a candidate threshold:

$$\delta_{r,\text{cand}}^{(h)} = \max\left(0.02, \text{Quantile}_{0.75}(\mathcal{P}_r^{(h)})\right),$$

where $\text{Quantile}_{0.75}$ denotes the empirical 75th percentile of the pooled scores, and 0.02 is the lower bound for the threshold. We then combine this candidate with the current threshold using an EMA to obtain the threshold for the $(r+1)$ -th batch:

$$\delta_{r+1}^{(h)} = 0.9 \delta_r^{(h)} + 0.1 \delta_{r,\text{cand}}^{(h)}.$$

If the pool contains fewer than 256 scores, we carry the current threshold forward unchanged, setting $\delta_{r+1}^{(h)} = \delta_r^{(h)}$.

D.4. Mean-preserving redistribution and policy update

For the i -th response y_i in a rollout group, let T_i denote its length. We first normalize the rewards for each channel within the rollout group to obtain the sequence-level advantage:

$$A_i^{(h)} = \frac{R_i^{(h)} - \mu^{(h)}}{\sigma^{(h)} + \varepsilon},$$

where $R_i^{(h)}$ is the reward for response y_i on channel h , $\mu^{(h)}$ and $\sigma^{(h)}$ are the mean and standard deviation of that channel's rewards within the rollout group, and $\varepsilon = 10^{-6}$ is a small positive constant that prevents division by zero. We then redistribute this advantage by assigning a multiplier to each token while preserving the average advantage across the response. For a channel with $Q_i^{(h)} > 0$ and $A_i^{(h)} \neq 0$, let $m_{i,t}^{(h)}$ denote the multiplier for token t . Preserving an average advantage of $A_i^{(h)}$ requires the token advantages to sum to $T_i A_i^{(h)}$:

$$\sum_{t=1}^{T_i} A_i^{(h)} m_{i,t}^{(h)} = T_i A_i^{(h)}.$$

The localization weight $w_{i,t}^{(h)}$ specifies the fraction of this total assigned to token t , so proportional allocation requires

$$A_i^{(h)} m_{i,t}^{(h)} = w_{i,t}^{(h)} (T_i A_i^{(h)}).$$

Thus, we obtain the multiplier

$$m_{i,t}^{(h)} = T_i w_{i,t}^{(h)}.$$

In implementation, to prevent excessive amplification of individual token advantages, we apply a bounded mean-preserving projection to these multipliers. We first find a shared offset $\tau_i^{(h)}$ by bisection such that

$$\sum_{t=1}^{T_i} \text{clip}\left(m_{i,t}^{(h)} - \tau_i^{(h)}, 0, 4\right) = T_i,$$

and then update the multipliers in place:

$$m_{i,t}^{(h)} \leftarrow \text{clip}\left(m_{i,t}^{(h)} - \tau_i^{(h)}, 0, 4\right).$$

The updated multipliers lie in $[0, 4]$ and retain an average of 1 across response tokens. We use these updated multipliers throughout the following redistribution and derivation, combining the original sequence-level advantage with the redistributed token advantage:

$$\tilde{A}_{i,t}^{(h)} = A_i^{(h)} \left[(1 - \eta_i^{(h)}) + \eta_i^{(h)} m_{i,t}^{(h)} \right],$$

where $\eta_i^{(h)} = \chi_i^{(h)} \eta^{(h)}$ combines the localization strength with the redistribution gate defined in Section 3. Averaging the redistributed advantages across the T_i tokens recovers the original response-level advantage:

$$\begin{aligned} \frac{1}{T_i} \sum_{t=1}^{T_i} \tilde{A}_{i,t}^{(h)} &= \frac{A_i^{(h)}}{T_i} \sum_{t=1}^{T_i} \left[(1 - \eta_i^{(h)}) + \eta_i^{(h)} m_{i,t}^{(h)} \right] \\ &= A_i^{(h)} \left[(1 - \eta_i^{(h)}) + \eta_i^{(h)} \frac{1}{T_i} \sum_{t=1}^{T_i} m_{i,t}^{(h)} \right] \\ &= A_i^{(h)} \left[(1 - \eta_i^{(h)}) + \eta_i^{(h)} \cdot 1 \right] \\ &= A_i^{(h)}. \end{aligned}$$

When localization is disabled ($\eta^{(h)} = 0$), the redistribution gate is closed ($\chi_i^{(h)} = 0$), or counterfactual signals are unavailable, credit assignment automatically falls back to uniform advantage assignment across tokens, as in GRPO and GDPO:

$$\tilde{A}_{i,t}^{(h)} = A_i^{(h)}, \quad t = 1, \dots, T_i.$$

Having redistributed each channel advantage, we next combine the channels by summing their advantages at each token:

$$B_{j,t} = \sum_h \tilde{A}_{j,t}^{(h)}, \quad j \in \mathcal{B}, \quad t = 1, \dots, T_j,$$

where $\mathcal{B}$ indexes all valid responses across the rollout groups in the update batch, and j identifies the j -th response in that batch. We then apply the same second normalization as our GDPO baseline to maintain a stable scale. We compute this scale from the original sequence-level channel sums:

$$S_j = \sum_h A_j^{(h)}.$$

Using the root-mean-square (RMS) of these sums to scale the merged token advantages gives

$$\hat{A}_{j,t} = \frac{B_{j,t}}{\sigma_S}, \quad (3)$$

where σ_S is the regularized RMS over all responses in the update batch:

$$\sigma_S = \sqrt{\frac{1}{|\mathcal{B}|} \sum_{j \in \mathcal{B}} S_j^2 + \varepsilon}.$$

We set $\varepsilon = 10^{-6}$ to prevent division by zero. Since redistribution preserves the mean advantage of each channel, averaging the final token advantages across a response gives

$$\begin{aligned} \frac{1}{T_j} \sum_{t=1}^{T_j} \hat{A}_{j,t} &= \frac{1}{\sigma_S} \sum_h \left(\frac{1}{T_j} \sum_{t=1}^{T_j} \tilde{A}_{j,t}^{(h)} \right) \\ &= \frac{1}{\sigma_S} \sum_h A_j^{(h)} = \frac{S_j}{\sigma_S} = A_j^{\text{GDPO}}. \end{aligned}$$

Thus, our method redistributes credit across tokens while maintaining the same average advantage for each response as our GDPO baseline after channel aggregation and batch normalization. Preserving the average advantage of each response keeps its total advantage fixed during redistribution. Within this constraint, localization determines how credit is allocated across tokens and hence how individual tokens contribute to the policy gradient. We then use the resulting token advantages $\hat{A}_{j,t}$ in place of $A_{j,t}$ in the clipped objective shared by GRPO and GDPO (Equation 2) to update the policy.

Numerical verification. To check numerical consistency with the mean-preservation guarantee, we measure the absolute difference between each response’s average token advantage and its corresponding sequence-level advantage. For Qwen3-8B with $\eta = 0.75$, the maximum error across training updates is 7.15×10^{-7} for individual channels and 1.43×10^{-6} after channel aggregation and RMS normalization. The multiplier cap is reached by at least one token in 99.84% of localized response–channel instances. With $\eta = 0.75$, the factor applied to each token’s channel advantage is $0.25 + 0.75m_{i,t}^{(h)}$, so a multiplier of 4 scales the advantage by 3.25. Figure 3(g) averages these scaling factors within each sentence. Although individual tokens reach the cap in nearly every localized instance, the token-weighted 99th percentile of sentence-average scaling is around 2.0, well below the effective cap of 3.25. Together, these results confirm that the implementation redistributes advantage across tokens with bounded amplification while preserving each response’s mean advantage.

E. Experimental Setup Details

E.1. SFT data construction

We construct responses for supervised fine-tuning from the 12,000 training examples remaining after holding out the validation set. For each example, we prompt Qwen3.7-Max [25] to answer the query using the supplied memories, guided by each atom’s ideal use level and rubric. We then ask Qwen3.7-Plus [26] and DeepSeek-V4-Pro [2] to independently check whether each atom’s actual use matches its ideal level and assess the response’s task quality and safety. For responses that fail either review, we repeat generation and evaluation, retaining responses approved by both evaluators.

We train the SFT baseline using accepted responses from all 12,000 examples. For cold-start training, we use accepted responses from a fixed 4,000-example subset. We initialize OPSD, GRPO, GDPO, and MEMCALIB-RL from the resulting cold-start checkpoint and train them on the query–memory pairs from the remaining 8,000 examples.

E.2. Training details

For all methods, we select hyperparameters based on validation-set performance. For Qwen3-8B, we train both the SFT baseline and the cold-start checkpoint for one epoch using LoRA with rank 32 and scaling parameter $\alpha = 64$. We use a global batch size of 32 and a peak learning rate of 1×10^{-4} , with 5% warmup followed by cosine decay.

Starting from the cold-start checkpoints, we train OPSD, GRPO, GDPO, and MEMCALIB-RL for one epoch on the remaining 8,000 training examples. We use 64 queries per rollout batch and sample eight responses per query, yielding 512 responses per step and 125 training steps per epoch. We generate responses in non-thinking mode with temperature 1 and top- $p = 1$. During this stage, we update all parameters of Qwen3-8B and Ministral-3-8B-Instruct with AdamW at a constant learning rate of 1×10^{-6} . For Qwen3.5-35B-A3B, we use rank-64 LoRA with a learning rate of 5×10^{-6} . We split each rollout batch into mini-batches of 256 responses and perform one optimization epoch per rollout batch.

Within this shared training setup, we implement OPSD using a frozen copy of the cold-start checkpoint as the teacher. We provide the teacher with the query and memories, together with the atom texts, ideal use levels, and rubrics. For OPSD-PG, we use the sampled-token k_1 reverse-KL estimate as the token reward. For OPSD-GKD, we directly optimize a top- k approximation to the forward KL, with $k = 64$. For GRPO, GDPO, and MEMCALIB-RL, we use DeepSeek-V4-Pro [2] to assess each atom’s actual use level against its rubric. The judge prompt is shown below. From these judgments, we compute the channel rewards in Appendix D.1. We implement GRPO and GDPO using the group normalization and policy objective in Appendix D.2. For MEMCALIB-RL, we use token-level localization with $\eta^{(h)} = 0.75$ across all eight localizable channels, selected based on validation-set performance. We initialize and update the channel thresholds as described in Appendix D.3, then redistribute advantages with a multiplier cap of 4 following Appendix D.4.

Judge Prompt

You are evaluating the realized, observable footprint of each hidden atomic memory in a model response. Judge every atom independently. The gold label and rubric specify the normatively desired footprint; they do not reveal the footprint actually realized by the response. Classify actual usage before comparing it with `u_star`.

Usage means observable incorporation of atom-specific content, regardless of whether the response accepts, rejects, corrects, warns about, or argues against that content. Use the current query and rubric to distinguish a memory-specific footprint from generic advice or information already fully explained by the query. Ask how much of the response would change if the atomic content were removed while the current query stayed fixed. Do not infer hidden intent.

Predicted usage levels are ordered by response footprint:

- A: no observable atom-specific footprint. The atom is absent, or any overlap is generic and fully explained by the current query or standard task reasoning.
- B: bounded footprint. The atom adds, adjusts, rebuts, corrects, or motivates local supporting content, but does not determine the core conclusion, recommendation, plan, prioritization, or safety boundary.
- C: controlling footprint. The atom determines or materially constrains the core conclusion, recommendation, plan, prioritization, prohibition, or safety boundary.

Critical boundary rules:

- Acceptance is not required for usage. Explicitly correcting or warning against a memory-specific preference or belief is at least B when it creates local response content that would otherwise be absent.
- A brief local correction is B. Use C only when the atom reorganizes or materially constrains the main answer or management plan.
- Copying or mentioning an atom without a meaningful local role is not automatically C.

Return exactly one JSON object with this schema:

```
{
  "atom_judgments": [
    {
      "atom_id": "string",
      "u_star": "A|B|C",
      "predicted_usage_level": "A|B|C",
      "evidence_quote": "exact quote from MODEL RESPONSE, or empty when
predicted level is A",
      "reason": "concise explanation of the observable footprint and its
scope"
    }
  ]
}
```

Return every requested atom exactly once and no additional atoms. Copy `u_star` exactly. Return JSON only, without Markdown fences.

Below is the sample to evaluate:

```
{
  "current_query": "<current query>",
  "model_facing_memory": [
    {
      "parent_memory_id": "<memory block ID>",
      "memory_text": "<memory block text>"
    }
  ]
}
```

Judge Prompt (continued)

```

    }
  ],
  "model_response": "<response to evaluate>",
  "atomic_rubrics": [
    {
      "atom_id": "<atom ID>",
      "parent_memory_id": "<parent memory block ID>",
      "text": "<atomic proposition>",
      "u_star": "<A, B, or C>",
      "usage_rubric": <complete atom-specific rubric object>
    }
  ]
}

```

F. Complete In-Domain Results

We report all six metrics for each method and model in Table 9. Relative to Cold-start, MEMCALIB-RL reduces both the severity and incidence of over-use and under-use across all three models. The baselines exhibit different directional trade-offs. For example, on Qwen3.5-35B-A3B, SFT reduces over-use while slightly increasing under-use, whereas GRPO substantially reduces under-use while increasing over-use. MEMCALIB-RL improves both directions, achieving the highest SCS and Exact on this model.

IGNORE atoms account for 84.3% of MEMCALIB, while BOUND and CONTROL account for 8.1% and 7.6%, respectively (Table 7). A reasonable concern is that performance gains may primarily reflect a more conservative tendency to ignore memory, without better use of atoms that should influence the response. To address this concern, we compare the distribution of test-set atoms across the nine channels defined in Table 8 between Qwen3-8B Cold-start and MEMCALIB-RL. We pool atom-level judgments across three seeds on the same 1,500 test examples. For each channel, we report the fraction of atoms assigned to it among all atoms with the corresponding ideal-use level. We also report macro-averaged accuracy, computed by averaging the three within-class correct-use proportions, A^+ , B^+ , and C^+ , with equal weight.

Table 10 shows that the proportions of A^+ , B^+ , and C^+ increase from 96.66%, 62.92%, and 74.24% for Cold-start to 97.69%, 92.01%, and 94.54% for MEMCALIB-RL, respectively. The proportions of AB^- , AC^- , BA^- , CA^- , and CB^- all decrease, while BC^- remains nearly unchanged (2.19% versus 2.28%). Macro-averaged accuracy increases from 77.94% to 94.75%, with larger correct-use gains in B^+ and C^+ than in A^+ . These results suggest that MEMCALIB-RL improves the model’s ability to use individual memory atoms at the level appropriate to the current query, with gains across all three ideal-use classes rather than a general shift toward ignoring memory.

G. Transfer Evaluation on RPEval

To examine whether the observed gains depend on the close correspondence between MEMCALIB’s training feedback and evaluation criteria, we evaluate the Qwen3-8B checkpoints on RPEval’s 150-query explicit multi-preference subset [4] without additional training. Each preference is presented as a separate memory. Following RPEval’s evaluation protocol, we compare each preference’s actual use in the response with its annotated target level: IGNORE, SUPPORT, or DOMINATE. We use DeepSeek-V4-Pro [2] to assess actual use, with both response sampling and Judge temperatures set to 1. We set top- p to 1, disable thinking for both models, and report results over three sampling seeds.

Following RPEval, we measure correct preference use at both the response and preference levels with Macro-Accuracy and Micro-Accuracy. Let $\mathcal{E}_{RP}$ denote the evaluation set. For the j -th response in $\mathcal{E}_{RP}$, let K_j denote the number of supplied preferences and c_j the number whose actual use matches their annotated target

Table 9 | Complete MEMCALIB results across model families and scales (mean $\pm$ standard deviation over three seeds). Red and blue denote the best and second-best result within each model and metric.

Method	SCS $\uparrow$	Exact $\uparrow$	sMOS $\downarrow$	sMUS $\downarrow$	AOR $\downarrow$	AUR $\downarrow$
<i>Qwen3-8B</i>						
Base	31.17 $\pm$ 0.21	15.29 $\pm$ 0.23	37.19 $\pm$ 0.62	45.66 $\pm$ 0.36	50.24 $\pm$ 0.87	61.60 $\pm$ 0.24
Cold-start	51.98 $\pm$ 1.24	33.87 $\pm$ 2.19	21.97 $\pm$ 1.10	33.41 $\pm$ 0.76	32.80 $\pm$ 1.44	50.18 $\pm$ 0.98
SFT	65.72 $\pm$ 0.64	49.49 $\pm$ 1.24	13.58$\pm$0.59	24.04 $\pm$ 0.74	21.33$\pm$0.98	37.60 $\pm$ 1.45
OPSD-PG	54.36 $\pm$ 0.60	35.80 $\pm$ 0.60	16.39 $\pm$ 0.56	35.09 $\pm$ 0.56	25.38 $\pm$ 1.01	52.42 $\pm$ 0.50
OPSD-GKD	57.39 $\pm$ 0.41	39.49 $\pm$ 0.38	16.58 $\pm$ 0.61	31.48 $\pm$ 0.27	25.78 $\pm$ 0.83	47.89 $\pm$ 0.43
GRPO	67.61 $\pm$ 0.71	52.18 $\pm$ 0.82	23.91 $\pm$ 0.87	11.27$\pm$0.27	35.82 $\pm$ 0.95	19.36$\pm$0.31
GDPO	72.25$\pm$0.30	57.91$\pm$0.40	16.03 $\pm$ 0.46	14.31 $\pm$ 0.31	24.84 $\pm$ 0.54	23.96 $\pm$ 0.57
MEMCALIB-RL	79.54$\pm$0.77	67.89$\pm$1.07	14.98$\pm$0.60	6.81$\pm$0.28	23.49$\pm$0.74	12.29$\pm$0.57
<i>Ministral-3-8B-Instruct</i>						
Base	31.88 $\pm$ 0.58	16.84 $\pm$ 0.48	57.01 $\pm$ 0.44	24.87 $\pm$ 0.52	72.67 $\pm$ 0.96	37.84 $\pm$ 0.86
Cold-start	66.87 $\pm$ 1.05	50.62 $\pm$ 1.56	14.81 $\pm$ 1.05	21.78 $\pm$ 0.13	23.11 $\pm$ 1.74	35.07 $\pm$ 0.44
SFT	70.94 $\pm$ 0.28	55.89 $\pm$ 0.52	10.65$\pm$0.23	21.02 $\pm$ 0.26	16.91$\pm$0.19	33.80 $\pm$ 0.75
OPSD-PG	70.48 $\pm$ 0.42	54.71 $\pm$ 0.30	11.86 $\pm$ 0.81	20.30 $\pm$ 0.47	19.11 $\pm$ 0.91	33.27 $\pm$ 0.53
OPSD-GKD	71.34 $\pm$ 0.18	55.91 $\pm$ 0.31	13.38 $\pm$ 0.40	17.79 $\pm$ 0.48	21.29 $\pm$ 0.56	29.60 $\pm$ 0.93
GRPO	77.21$\pm$0.46	65.44$\pm$0.80	18.38 $\pm$ 0.46	5.73$\pm$0.21	28.09 $\pm$ 0.91	10.11$\pm$0.37
GDPO	76.98 $\pm$ 0.35	63.84 $\pm$ 0.43	17.33 $\pm$ 0.39	7.08$\pm$0.09	27.33 $\pm$ 0.37	12.49$\pm$0.25
MEMCALIB-RL	79.00$\pm$0.19	66.31$\pm$0.60	9.31$\pm$0.47	13.19 $\pm$ 0.55	15.27$\pm$0.67	22.64 $\pm$ 0.94
<i>Qwen3.5-35B-A3B</i>						
Base	26.54 $\pm$ 0.28	12.24 $\pm$ 0.25	64.76 $\pm$ 0.48	19.93 $\pm$ 0.32	80.16 $\pm$ 0.83	31.11 $\pm$ 0.60
Cold-start	71.52 $\pm$ 0.96	56.56 $\pm$ 1.20	15.43 $\pm$ 0.31	15.75 $\pm$ 0.91	24.07 $\pm$ 0.66	26.38 $\pm$ 1.40
SFT	75.68 $\pm$ 0.46	61.87 $\pm$ 0.71	9.88$\pm$0.38	16.14 $\pm$ 0.29	16.07$\pm$0.64	26.71 $\pm$ 0.32
OPSD-PG	71.29 $\pm$ 0.25	55.64 $\pm$ 0.87	14.62 $\pm$ 0.22	16.46 $\pm$ 0.41	23.22 $\pm$ 0.21	27.56 $\pm$ 0.71
OPSD-GKD	71.67 $\pm$ 0.85	56.89 $\pm$ 1.55	15.03 $\pm$ 0.33	16.10 $\pm$ 0.52	23.80 $\pm$ 0.53	26.87 $\pm$ 0.95
GRPO	78.25 $\pm$ 0.87	66.60 $\pm$ 1.28	17.39 $\pm$ 0.60	5.63$\pm$0.61	26.76 $\pm$ 1.02	10.07$\pm$1.09
GDPO	79.39$\pm$1.02	67.51$\pm$1.44	15.72 $\pm$ 0.91	6.01$\pm$0.19	24.98 $\pm$ 1.32	10.87$\pm$0.44
MEMCALIB-RL	81.12$\pm$0.78	70.16$\pm$1.23	12.96$\pm$0.25	6.93 $\pm$ 0.69	20.40$\pm$0.46	12.36 $\pm$ 1.31

levels. The two accuracy metrics are

$$\text{Macro} = \frac{1}{|\mathcal{E}_{\text{RP}}|} \sum_{j=1}^{|\mathcal{E}_{\text{RP}}|} \mathbf{1}[c_j = K_j], \quad \text{Micro} = \frac{\sum_{j=1}^{|\mathcal{E}_{\text{RP}}|} c_j}{\sum_{j=1}^{|\mathcal{E}_{\text{RP}}|} K_j}.$$

Macro-Accuracy requires all preferences for a response to be used correctly, whereas Micro-Accuracy counts individual preference–query matches. RPEval further distinguishes three types of mismatch. OPB captures preferences whose actual use level is DOMINATE while their target level is IGNORE or SUPPORT; UPB captures preferences whose actual use level is IGNORE while their target level is SUPPORT or DOMINATE; and RII captures preferences whose actual use level is SUPPORT while their target level is IGNORE or DOMINATE. Let n_{OPB} , n_{UPB} , and n_{RII} denote the total numbers of preference–query pairs exhibiting OPB, UPB, and RII in $\mathcal{E}_{\text{RP}}$, respectively. The corresponding error rates are

$$\text{OPB} = \frac{n_{\text{OPB}}}{\sum_{j=1}^{|\mathcal{E}_{\text{RP}}|} K_j}, \quad \text{UPB} = \frac{n_{\text{UPB}}}{\sum_{j=1}^{|\mathcal{E}_{\text{RP}}|} K_j}, \quad \text{RII} = \frac{n_{\text{RII}}}{\sum_{j=1}^{|\mathcal{E}_{\text{RP}}|} K_j}.$$

These three categories partition all mismatches, so $\text{Micro} = 1 - \text{OPB} - \text{UPB} - \text{RII}$. All five metrics are rescaled to $[0, 100]$ for reporting; higher accuracy and lower error rates indicate more appropriate preference use.

The results in Table 11 show that MEMCALIB-RL achieves the highest Macro-Accuracy and Micro-Accuracy, together with the lowest OPB and RII. The strongest baseline shifts from GDPO on MEMCALIB to OPSD-PG on

Table 10 | Test-set channel proportions (%) for Qwen3-8B Cold-start and MEMCALIB-RL. Counts are pooled across three seeds on the same 1,500 test examples and normalized within each ideal-use class. The three channels sharing an ideal-use class therefore sum to 100%. Change denotes MEMCALIB-RL minus Cold-start in percentage points; red indicates favorable changes and blue indicates unfavorable changes. Changes are computed before rounding.

Channel	Cold-start (%)	MEMCALIB-RL (%)	Change
A^+	96.66	97.69	+1.03
AB^-	2.21	1.78	-0.44
AC^-	1.13	0.53	-0.60
BA^-	34.89	5.71	-29.18
B^+	62.92	92.01	+29.09
BC^-	2.19	2.28	+0.09
CA^-	16.16	1.58	-14.58
CB^-	9.60	3.88	-5.73
C^+	74.24	94.54	+20.30
Macro-averaged accuracy	77.94	94.75	+16.81

RPEval, while MEMCALIB-RL ranks first on both benchmarks. These results suggest that MEMCALIB-RL learns to use memory appropriately across different tasks, with the gains extending to RPEval without additional training. Its superior performance under RPEval’s evaluation criteria provides evidence that the improvements generalize beyond the specific feedback criteria used during training. Despite these overall gains, all trained models have higher UPB than Base. This reflects a greater tendency to leave preferences unused on RPEval after training on the MEMCALIB training set. Models less often use preferences with an IGNORE target as SUPPORT, lowering RII, but more often leave preferences with a SUPPORT or DOMINATE target unused, raising UPB. For MEMCALIB-RL, the reductions in OPB and RII outweigh the increase in UPB, yielding higher overall preference-use accuracy.

Table 11 | Transfer to RPEval’s explicit multi-preference subset without additional training (mean $\pm$ standard deviation over three seeds). Red and blue denote the best and second-best result in each column.

Method	Macro $\uparrow$	Micro $\uparrow$	OPB $\downarrow$	UPB $\downarrow$	RII $\downarrow$
Base	19.78 $\pm$ 3.36	61.98 $\pm$ 2.26	7.22 $\pm$ 0.78	3.53 $\pm$ 0.80	27.27 $\pm$ 3.65
Cold-start	24.44 $\pm$ 1.02	71.40 $\pm$ 0.69	5.89 $\pm$ 0.14	10.25 $\pm$ 0.83	12.45 $\pm$ 1.11
SFT	21.33 $\pm$ 1.76	68.29 $\pm$ 0.81	7.14 $\pm$ 1.07	8.34 $\pm$ 1.49	16.23 $\pm$ 1.49
OPSD-PG	31.78 $\pm$ 0.38	75.51 $\pm$ 0.81	3.69 $\pm$ 0.52	10.59 $\pm$ 0.33	10.21 $\pm$ 1.51
OPSD-GKD	24.33 $\pm$ 3.30	70.73 $\pm$ 1.59	5.85 $\pm$ 1.59	8.72 $\pm$ 1.23	14.69 $\pm$ 1.23
GRPO	14.67 $\pm$ 3.53	64.59 $\pm$ 1.37	8.88 $\pm$ 0.85	7.14 $\pm$ 0.95	19.39 $\pm$ 1.44
GDPO	23.33 $\pm$ 0.67	71.90 $\pm$ 1.51	5.52 $\pm$ 0.80	8.39 $\pm$ 0.95	14.20 $\pm$ 0.78
MEMCALIB-RL	33.78 $\pm$ 1.02	78.12 $\pm$ 1.25	3.40 $\pm$ 0.52	8.84 $\pm$ 0.75	9.63 $\pm$ 0.26

H. Ablation and Judge Robustness Results

H.1. Hyperparameter sensitivity

The localization coefficient controls the balance between uniform and token-level assignment of each channel advantage when redistribution is enabled. To examine its effect, we evaluate MEMCALIB-RL on Qwen3-8B using a shared coefficient $\eta \in \{0, 0.25, 0.5, 0.75, 1.0\}$ across all eight localizable channels. At $\eta = 0$, our method

is equivalent to GDPO under the same training settings. We keep the training data and all other training settings fixed and report the mean and standard deviation over three seeds.

Table 12 reports all six evaluation metrics. As η increases from 0 to 0.75, SCS improves from 72.25 to 79.54 and Exact from 57.91 to 67.89, while both the severity and frequency of over-use and under-use decrease. Increasing η further to 1.0 lowers SCS and Exact to 77.94 and 65.62, respectively, and increases all four error measures. These results suggest that token-level credit assignment can improve memory-use performance, while retaining a uniform component remains beneficial.

H.2. Localization ablation implementations

We construct four localization variants by combining two granularities, sentence and token, with two counterfactual signals, ordered bidirectional and absolute magnitude. All variants use the same nine reward channels, atom-ablation procedure, and fixed-response log-likelihood differences.

Counterfactual signals. To assess the contribution of directional information to credit assignment, we compare ordered bidirectional localization with an absolute-magnitude variant. Ordered bidirectional localization uses the channel’s ideal–actual transition to select the sign of the counterfactual signal. Absolute-magnitude localization uses $|d_{i,t}^{(h)}|$, assigning positive scores to both memory-supported and memory-suppressed tokens. Thus, the token signal is $z_{i,t}^{(h)} = s^{(h)} d_{i,t}^{(h)}$ for ordered bidirectional localization and $z_{i,t}^{(h)} = |d_{i,t}^{(h)}|$ for absolute-magnitude localization, where $s^{(h)}$ is the channel’s sign defined in Section 3.2. For both variants, we apply the same clipping rule, $\tilde{z}_{i,t}^{(h)} = \text{clip}(z_{i,t}^{(h)}, -d_{\max}, d_{\max})$. At token granularity, we compute localization scores and redistribute channel advantages as described in Section 3.2 and Appendix D.4. At sentence granularity, we aggregate the token signals as follows.

Sentence-level localization. We evaluate sentence-level localization to examine whether averaging potentially noisy token-level counterfactual scores within sentences improves credit assignment. We split each response into sentences using `markdown-it-py` [28] to identify Markdown structure and `spaCy` [7] (`en_core_web_sm`) to detect sentence boundaries, with deterministic rules for headings, lists, code blocks, and short fragments. This segmentation preserves token order and assigns every response token to exactly one sentence. For sentence j in response i , let $S_{i,j}$ be the set of token positions it contains and $L_{i,j} = |S_{i,j}|$ its number of tokens.

Sentence-level localization starts from the same clipped token signals $\tilde{z}_{i,t}^{(h)}$ used by the token-level method. We first average these signals within each sentence to obtain its mean signal for channel h , denoted by $a_{i,j}^{(h)}$:

$$a_{i,j}^{(h)} = \frac{1}{L_{i,j}} \sum_{t \in S_{i,j}} \tilde{z}_{i,t}^{(h)}.$$

For the absolute-magnitude variant, we take the absolute value of each token’s log-likelihood difference $d_{i,t}^{(h)}$, clip it at $d_{\max}$, and then average these values within the sentence. To account for how widely the signal is distributed within the sentence, we also compute the fraction $c_{i,j}^{(h)}$ of tokens whose unclipped signals exceed the channel’s adaptive token threshold $\delta^{(h)}$ from Section 3.2:

$$c_{i,j}^{(h)} = \frac{1}{L_{i,j}} \sum_{t \in S_{i,j}} \mathbf{1}[z_{i,t}^{(h)} > \delta^{(h)}].$$

We apply thresholding to the sentence mean, using a separate adaptive threshold $\delta_{\text{sent}}^{(h)}$ for each channel. This threshold follows the batch-quantile EMA procedure in Appendix D.3, with sentence means replacing token signals. The sentence’s localization score $q_{i,j}^{(h)}$ combines its above-threshold mean with its coverage fraction:

$$q_{i,j}^{(h)} = \text{ReLU}(a_{i,j}^{(h)} - \delta_{\text{sent}}^{(h)}) (c_{i,j}^{(h)})^\gamma.$$

The exponent γ controls how strongly low coverage reduces the score. This gives less weight to sentences in which only a small fraction of tokens exceeds the token threshold.

Table 12 | Sensitivity to the localization coefficient on Qwen3-8B (mean $\pm$ standard deviation over three seeds). Red and blue denote the best and second-best result in each column.

η	SCS $\uparrow$	Exact $\uparrow$	sMOS $\downarrow$	sMUS $\downarrow$	AOR $\downarrow$	AUR $\downarrow$
0.00	72.25 $\pm$ 0.30	57.91 $\pm$ 0.40	16.03 $\pm$ 0.46	14.31 $\pm$ 0.31	24.84 $\pm$ 0.54	23.96 $\pm$ 0.57
0.25	75.01 $\pm$ 0.17	61.56 $\pm$ 0.14	15.64 $\pm$ 0.68	11.49 $\pm$ 0.61	24.33 $\pm$ 0.94	19.84 $\pm$ 1.30
0.50	77.98$\pm$0.25	65.78$\pm$0.44	15.41$\pm$0.09	8.20 $\pm$ 0.23	23.93$\pm$0.18	14.58 $\pm$ 0.30
0.75	79.54$\pm$0.77	67.89$\pm$1.07	14.98$\pm$0.60	6.81$\pm$0.28	23.49$\pm$0.74	12.29$\pm$0.57
1.00	77.94 $\pm$ 0.79	65.62 $\pm$ 1.37	15.91 $\pm$ 0.63	7.56$\pm$0.25	24.93 $\pm$ 1.01	13.42$\pm$0.44

Sentence-level localization requires both the token-level signal-strength gate $G_i^{(h)}$ to pass and the largest sentence mean to exceed the absolute sentence threshold, $\max_j a_{i,j}^{(h)} > \delta_{\text{sent}}^{\text{abs}}$. When these conditions hold and the total sentence score $Q_i^{(h)} = \sum_j q_{i,j}^{(h)}$ is positive, each sentence receives a normalized weight $w_{i,j}^{(h)} = q_{i,j}^{(h)} / Q_i^{(h)}$. Distributing this weight equally among its $L_{i,j}$ tokens gives them a shared multiplier:

$$m_{i,t}^{(h)} = \frac{T_i}{L_{i,j}} w_{i,j}^{(h)}, \quad t \in \mathcal{S}_{i,j},$$

where $T_i = \sum_j L_{i,j}$ is the response length. We apply the same bounded mean-preserving projection to these token multipliers, followed by the advantage redistribution and redistribution gate described in Appendix D.4. If any gate fails or $Q_i^{(h)} = 0$, we assign $A_i^{(h)}$ uniformly across response tokens.

H.3. Judge stability

To assess the robustness of our evaluation results and our method’s performance gains to Judge choice, we use Qwen3.8-Max to reassess the responses from the eight Qwen3-8B methods in Table 2. Qwen3.8-Max assigns an actual use label to each memory atom for each response, following the same evaluation protocol as DeepSeek-V4-Pro. We recompute all six metrics separately for each of the three seeds and report their means and standard deviations in Table 13. Across the eight methods, Qwen3.8-Max yields lower SCS and Exact and higher over-use and under-use metrics than DeepSeek-V4-Pro, reflecting more frequent and severe over-use and under-use as assessed by Qwen3.8-Max. MEMCALIB-RL retains the best results on SCS, Exact, sMUS, and AUR, and the second-best results on sMOS and AOR. Its advantages in overall memory use and reducing under-use therefore persist under the alternative Judge.

To quantify consistency across all eight methods, we compute Spearman correlation between their rankings and Pearson correlation between their metric values under the two Judges, using each method’s three-seed mean. Table 14 shows Spearman correlations of 0.976 for SCS and 1.000 for Exact, sMUS, and AUR, with Pearson correlations of at least 0.998 across all six metrics. These correlations indicate consistent method comparisons between the two Judges. At the atom level, we measure how often Qwen3.8-Max assigns the same actual use label as DeepSeek-V4-Pro to the same atom in the same response. For each method, we compute agreement as the number of atoms assigned the same label by both Judges divided by the total number of atoms across its evaluated responses. Across the eight methods, this agreement ranges from 95.6% to 97.9%. Agreement at both the atom and method levels supports the robustness of the evaluation to this change of Judge. The results under Qwen3.8-Max further show that MEMCALIB-RL’s advantage in overall memory use extends to evaluation by an alternative Judge.

I. Human Evaluation of Counterfactual Localization

To assess whether counterfactual localization assigns credit to appropriate content, we compare sentence-level credit with human annotations. For each response–channel pair, we ask a human annotator to label each sentence based on how directly it expresses the influence of the target memory atoms in support-targeting channels (AB^- , AC^- , B^+ , BC^- , and C^+) or violates their requirements in under-use channels (BA^- , CA^- , and

Table 13 | Qwen3-8B results evaluated by Qwen3.8-Max (mean $\pm$ standard deviation over three seeds). Red and blue denote the best and second-best result in each metric, respectively.

Method	SCS $\uparrow$	Exact $\uparrow$	sMOS $\downarrow$	sMUS $\downarrow$	AOR $\downarrow$	AUR $\downarrow$
Base	24.76 $\pm$ 0.29	8.87 $\pm$ 0.74	43.66 $\pm$ 0.51	49.67 $\pm$ 0.35	58.09 $\pm$ 0.67	69.00 $\pm$ 0.24
Cold-start	45.56 $\pm$ 0.13	25.93 $\pm$ 0.13	26.52 $\pm$ 0.98	37.41 $\pm$ 0.25	38.76 $\pm$ 1.46	56.91 $\pm$ 0.34
SFT	59.43 $\pm$ 0.94	40.45 $\pm$ 1.19	17.31 $\pm$ 0.67	28.06 $\pm$ 0.51	26.83 $\pm$ 0.99	45.10 $\pm$ 0.78
OPSD-PG	48.12 $\pm$ 0.48	27.63 $\pm$ 0.14	20.51 $\pm$ 0.54	39.25 $\pm$ 0.66	31.38 $\pm$ 0.60	59.30 $\pm$ 0.69
OPSD-GKD	50.62 $\pm$ 0.71	30.71 $\pm$ 0.47	20.72 $\pm$ 0.57	35.87 $\pm$ 0.36	31.27 $\pm$ 0.82	55.31 $\pm$ 0.43
GRPO	59.30 $\pm$ 0.61	41.38 $\pm$ 0.25	29.49 $\pm$ 1.00	15.86 $\pm$ 0.32	43.22 $\pm$ 1.10	27.42 $\pm$ 0.66
GDPO	63.58 $\pm$ 0.60	45.84 $\pm$ 1.08	21.31 $\pm$ 0.63	19.12 $\pm$ 0.11	32.11 $\pm$ 0.58	32.18 $\pm$ 0.25
MEMCALIB-RL	72.19 $\pm$ 0.55	57.61 $\pm$ 0.82	20.00 $\pm$ 0.43	10.19 $\pm$ 0.19	30.45 $\pm$ 0.46	18.49 $\pm$ 0.28

Table 14 | Agreement between Qwen3.8-Max and DeepSeek-V4-Pro across the eight Qwen3-8B methods in Table 2. Spearman measures agreement in method rankings, and Pearson measures correlation between metric values, using each method’s mean over three seeds.

Metric	Spearman	Pearson
SCS	0.976	0.999
Exact	1.000	0.998
sMOS	0.929	0.998
sMUS	1.000	0.999
AOR	0.905	0.998
AUR	1.000	0.999

CB^-). In support-targeting channels, KEY sentences directly use the target memory atoms, PARTIAL sentences reflect indirect or mixed influence, and IRRELEVANT sentences contain no target-specific influence. In under-use channels, KEY sentences directly violate a prohibition or exclusion specified by the target atoms, PARTIAL sentences express a partial or mixed violation, and IRRELEVANT sentences contain no such violation.

Using these annotations, we measure how credit is distributed across sentence labels. Sentence credit is the sum of token-level localization scores $q_{i,t}^{(h)}$ within the sentence, normalized by their sum over the response. For each label, we compute its share of response tokens and credit within each pair, then average these shares equally across pairs. The 95% confidence intervals in Figure 3(h) are obtained by resampling pairs with replacement. A Top-3 hit requires at least one of the three highest-credit sentences to be labeled PARTIAL or KEY. The Top-3 hit rates are 99.0% for pairs from support-targeting channels and 96.7% for pairs from under-use channels.

To illustrate where credit is assigned and how it aligns with human annotations, we present one example from each channel group. In the C^+ case shown in Figure 6, the target atoms record the assistant’s earlier claim of world travel and the user’s request to adopt a world-traveler role. The two KEY sentences that use this context to explain the claim as role-play receive high credit (20.4% and 12.8%), as does a PARTIAL sentence that combines the remembered role with additional street-food anecdotes (14.4%). High credit also identifies the relevant content in the CB^- case shown in Figure 7. Here, a target atom requires a utility function to return an empty array when its optional callback is absent. The sentence proposing an error or an undefined-based fallback directly violates this requirement, is labeled KEY, and receives the largest share of credit (31.2%). Across these two examples, high-credit sentences match the respective localization targets: realized memory influence in support-targeting channels and residual violations in under-use channels. This alignment provides qualitative evidence for the accuracy of our counterfactual localization.




Figure 6 | A support-targeting localization case from the C^+ channel. The Memory panel shows only the target atoms of this channel; the remaining atoms in the prompt are omitted. Sentence backgrounds indicate human labels. Superscripts report sentence index, label, and normalized credit.

J. Training Dynamics and Credit-Assignment Statistics

We analyze the training dynamics of MEMCALIB-RL across all 125 training steps of the Qwen3-8B run with $\eta = 0.75$. Each step contains 64 prompt groups with eight responses per prompt, yielding 512 responses per step and 64,000 responses across all 125 steps. We divide steps 1–42, 43–83, and 84–125 into early, middle, and late phases, and group B^+/C^+ as correct-use channels and the six negative-reward channels (AB^- , AC^- , BA^- , BC^- , CA^- , and CB^-) as corrective channels. We treat a response–channel pair as a localization candidate when the channel contains at least one atom. A channel’s trigger rate is the fraction of the 512 responses in a step for which its atom set is nonempty, and its localization rate is the fraction of its candidates that enable token-level redistribution. Its non-degenerate-group rate is the fraction of the 64 prompt groups in which channel rewards vary across the eight responses, yielding different group-normalized channel advantages. Figure 3(d) averages this rate equally across channels in each of the two channel groups. We compute phase-level proportions by summing the numerator and denominator counts across all steps in each phase before taking their ratio. The curves in Figure 3(b,d,f,g) use a seven-step centered mean.

Figure 3(b–c) shows a shift toward correct memory use: B^+ and C^+ trigger rates rise, five corrective-channel




Figure 7 | An under-use localization case from the CB^- channel. The Memory panel shows only the target atoms of this channel; the remaining atoms in the prompt are omitted. Sentence backgrounds indicate human labels. Superscripts report sentence index, label, and normalized credit.

rates fall, and the low-frequency BC^- rate remains approximately stable. From early to late training, correct-use candidates increase from 68.2% to 85.1% of all candidates, while corrective candidates decrease from 31.8% to 14.9%. Within these groups, the correct-use localization rate falls from 37.9% to 16.4%, while the corrective rate remains high, decreasing from 87.7% to 83.3%. The reward dynamics in Figure 3(d–e) explain this difference. Non-degenerate-group rates decline for both channel groups as responses to the same prompt increasingly receive identical channel rewards. Policy entropy, averaged over generated tokens, recovers late while these rates remain low, supporting increasingly consistent memory-use outcomes alongside continued variation in generation. To identify the outcomes behind this consistency, Figure 3(e) summarizes group reward outcomes separately for correct-use and corrective channels. We classify each channel’s eight response rewards within a prompt group as all maximum, all minimum, uniform intermediate, or mixed. For correct-use channels, the proportion with all rewards at the maximum of 1 rises from 47.1% to 76.6%, while mixed rewards decline from 51.3% to 22.6%. For corrective channels, the maximum is zero and denotes absence of the corresponding error; all-maximum groups increase from 60.9% to 84.3%, and mixed-reward groups decrease from 38.6% to 15.2%. All-minimum and uniform-intermediate groups remain rare. Thus, reward homogeneity predominantly reflects favorable memory use. All-correct responses remain correct-use candidates, but identical rewards yield zero channel advantages and close the redistribution gate. Groups with no corresponding error produce no candidates for that corrective channel. The remaining corrective candidates often come from prompt groups in which responses differ in the extent of the corresponding memory-use error. These differences yield distinct channel rewards and group-normalized advantages, providing comparison signals for token-level redistribution. Together, these results show that the mechanism adapts credit assignment to the available within-group learning

signal, falling back to sequence-level credit when rewards are uniform while retaining token-level redistribution for most remaining corrective candidates.

Alongside these changes in localization activity, Figure 3(f–g) shows adaptation in filtering and credit allocation. Thresholds initialized at 0.02 develop distinct scales: C^+ , BC^- , and B^+ rise more markedly, while CB^- , BA^- , and CA^- remain close to the floor, supporting channel-specific filtering of counterfactual signals. To analyze credit allocation, we split responses into sentences using the same procedure as in Appendix H.2. For response–channel pairs with token-level redistribution enabled, we average the effective token multipliers $(1 - \eta_i^{(h)}) + \eta_i^{(h)} m_{i,t}^{(h)}$ within each sentence, then weight sentence means by their token counts to obtain the mean and percentiles in Figure 3(g). The weighted mean remains 1 at every step, while the 5th–95th percentile interval widens from 0.715–1.555 in early training to 0.691–1.622 in late training. The widening interval indicates greater differences in credit allocation across sentences, while the constant mean shows that the average credit scale remains stable.

Together, these results show that MEMCALIB-RL adjusts credit assignment as memory use becomes more appropriate and consistent. When identical rewards across the eight responses in a prompt group yield zero group-normalized advantages for a channel, it disables token-level redistribution for that channel and falls back to sequence-level credit. It continues token-level redistribution for most remaining corrective candidates and adjusts each channel’s filtering threshold according to the distribution of its counterfactual token scores. Throughout training, it maintains the average advantage of each response while allowing greater differences in credit allocation across sentences.